Deploy Workloads with Lakeflow Jobs

Databricks Academy

September 2025

Agenda

Course Sections
• Introduction to Data Engineering in Databricks
• Lakeflow Jobs Core Concepts
• Creating and Scheduling Jobs
• Advanced Lakeflow Jobs Features

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Course Learning Objectives

- Understand the role of Lakeflow Jobs within the Databricks ecosystem as a unified orchestration platform for data, analytics, and AI workloads.
- Design and implement data workloads using Directed Acyclic Graphs (DAGs), demonstrating the relationship between jobs, tasks, and dependencies.
- Configure various job scheduling options including manual, scheduled, file arrival, and continuous triggers to automate workflow execution.
- Implement advanced workflow features such as conditional task execution, run-if dependencies, for each loop and repair runs to create robust, fault-tolerant data pipelines.
- Apply best practices for production workloads including selecting appropriate compute options, implementing modular orchestration, and utilizing proper error handling techniques.

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Course Prerequisites (REQUIRED)

Fundamental Knowledge of the Databricks Platform

- Databricks Workspaces
- Apache Spark
- Delta Lake and the Medallion Architecture
- Unity Catalog

Intermediate Python Knowledge

- Creating functions and working with variables
- Reading data into DataFrames
- Applying basic transformations (select, filter, groupBy, etc.)

Experience Transforming Data with SQL

- Writing intermediate-level queries
- Creating and manipulating columns
- Aggregating data
- Filtering results

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lab Exercise Environment

Technical Details

- Your lab environment is provided by Vocareum.
- It will open in a new tab.
- It has been configured with the permissions and resources required to accomplish the tasks outlined in the lab exercise.
- Third party cookies must be enabled in your browser for Vocareum user experience to work properly.
- Make sure to enable pop ups!

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Introduction to Data Engineering in Databricks

Deploy Workloads with Lakeflow Jobs

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Section Learning Objectives

- Understand the key components of data engineering on the Databricks platform, including Connect, Spark Declarative Pipelines and Jobs.
- Explain what Lakeflow Jobs are and their main benefits for unified orchestration of data, analytics, and AI workloads.
- Identify the core capabilities and use cases enabled by Lakeflow Jobs within the Databricks ecosystem.

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Agenda

Introduction to Data Engineering in Databricks
• Data Engineering in Databricks
• What is Lakeflow Jobs?

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Introduction to Data Engineering in Databricks

LECTURE

Data Engineering in Databricks

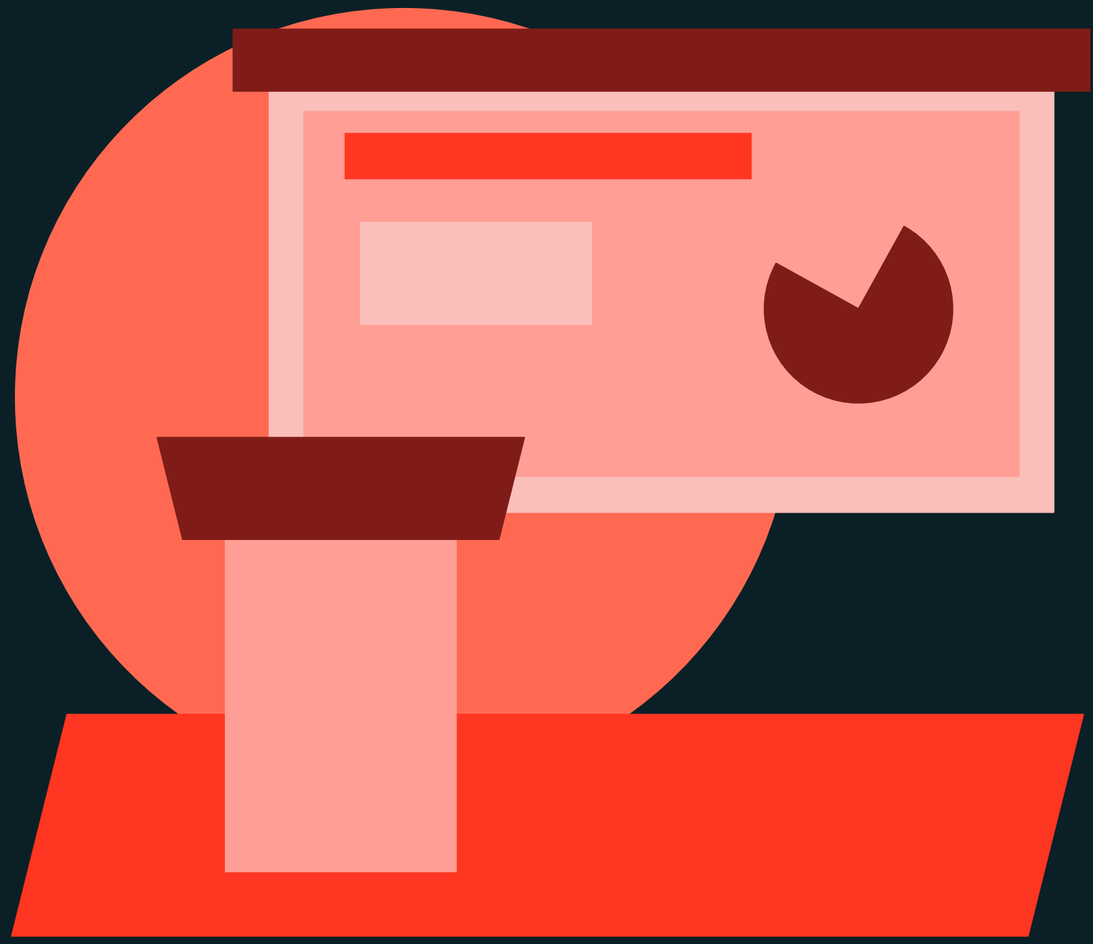

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Data Engineering in Databricks

databricks LAKEFLOW

UNIFIED DATA ENGINEERING FOR THE DATA INTELLIGENCE PLATFORM

CONNECT

Efficient ingestion connectors

SPARK DECLARATIVE PIPELINES (Previously DLT)

Accelerated ETL development

JOBS (Previously Workflows)

Reliable orchestration for analytics and AI

INDUSTRY LEADING DATA PROCESSING ENGINE
(Apache Spark + Structured Streaming)

UNIFIED GOVERNANCE	Unity Catalog
OPTIMIZED STORAGE	DELTA LAKE, Parquet, ICEBERG

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Data Engineering in Databricks

This course will focus on **orchestrating every workload** (Notebooks, SQL, dashboards, Pipelines, and more)

databricks LAKEFLOW

UNIFIED DATA ENGINEERING FOR THE DATA INTELLIGENCE PLATFORM

CONNECT

Efficient ingestion connectors

SPARK DECLARATIVE PIPELINES (Previously DLT)

Accelerated ETL development

JOBS (Previously Workflows)

Reliable orchestration for analytics and AI

INDUSTRY LEADING DATA PROCESSING ENGINE
(Apache Spark + Structured Streaming)

UNIFIED GOVERNANCE	Unity Catalog
OPTIMIZED STORAGE	DELTA LAKE, Parquet, ICEBERG

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Introduction to Data Engineering in Databricks
LECTURE

What is Lakeflow Jobs?

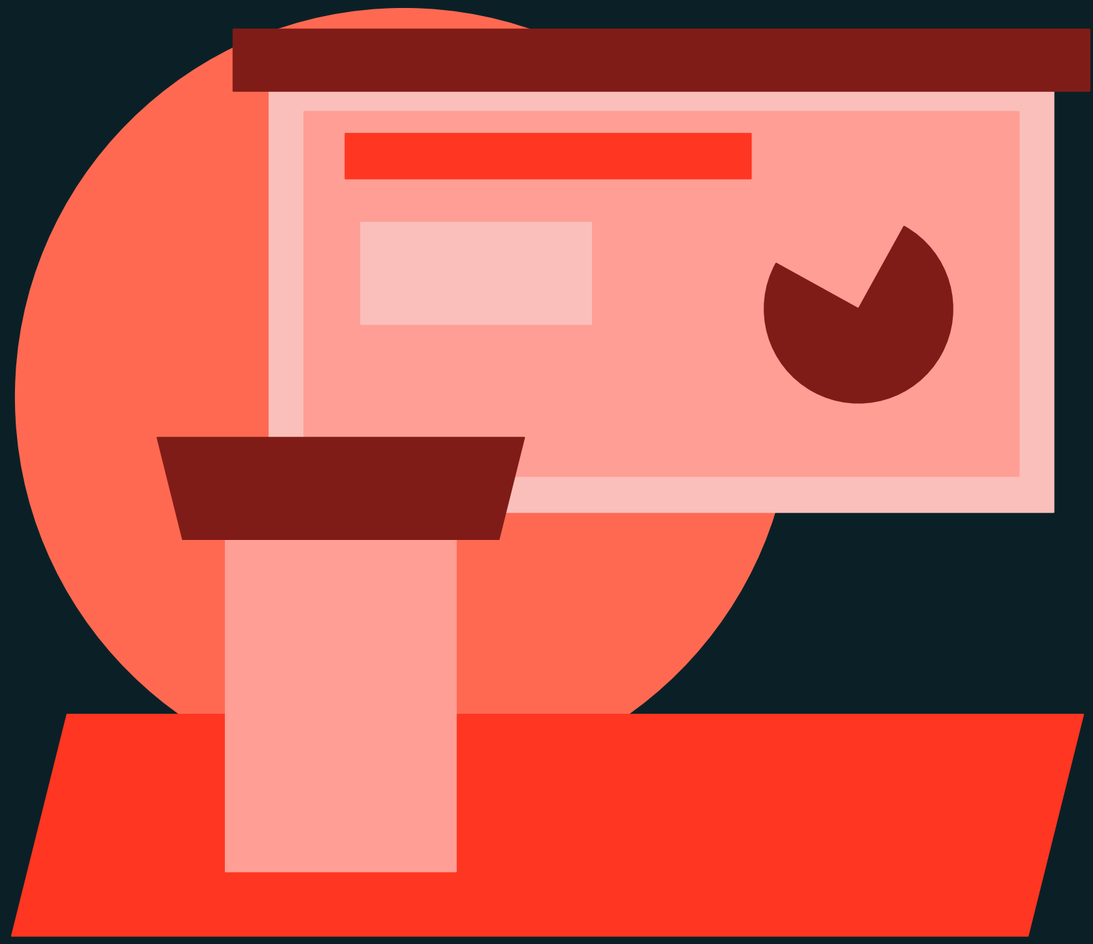

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

There are many ways to orchestrate your workloads

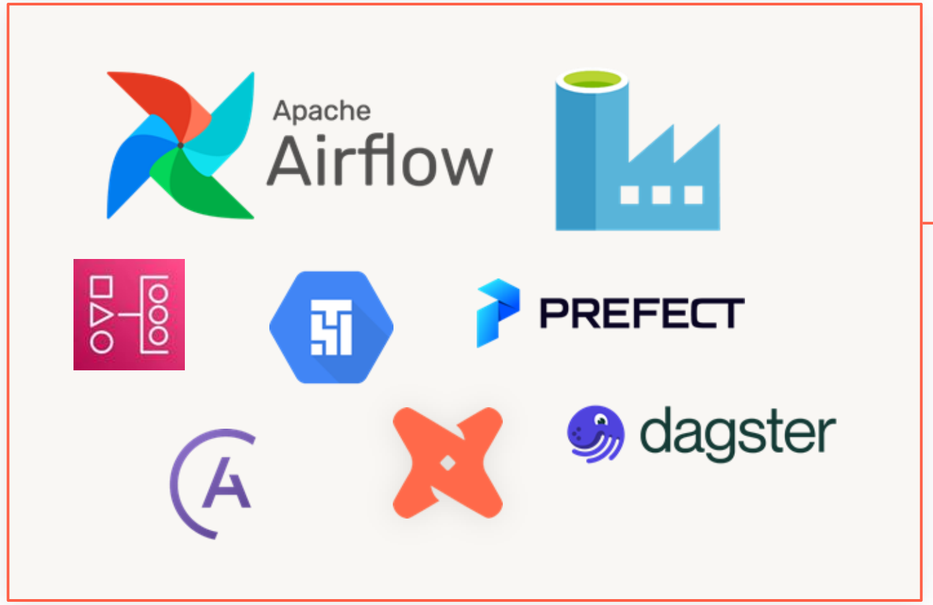

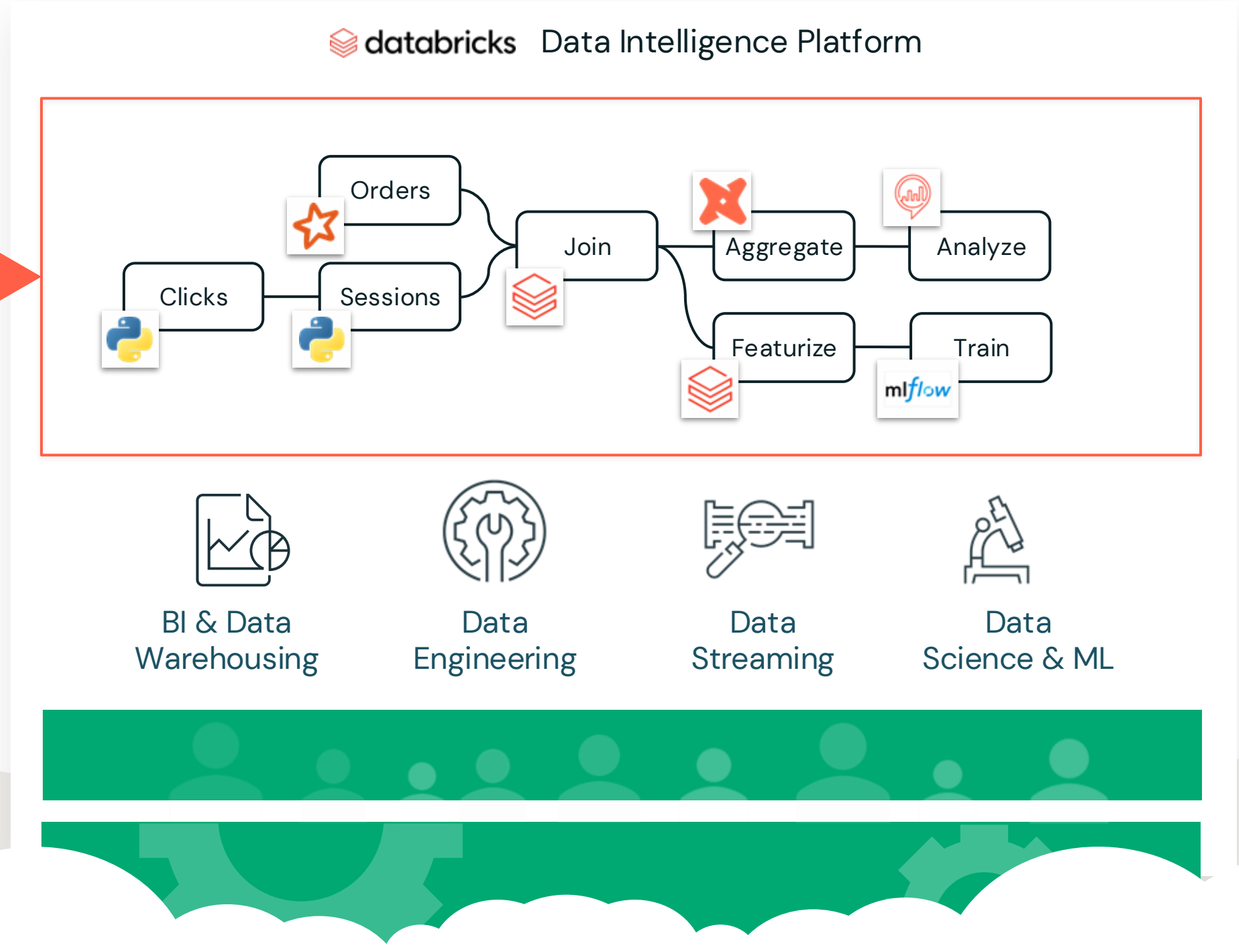

External orchestrators create challenges

Hard to use for many practitioners	Difficult to understand root cause when issues occur	Complex architecture to manage and maintain
Data teams are less productive	Bad data lowers value of downstream applications	Higher cost of ownership and lower reliability

These tools are <u>not unified</u> with your Lakehouse

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

What is Lakeflow Jobs?

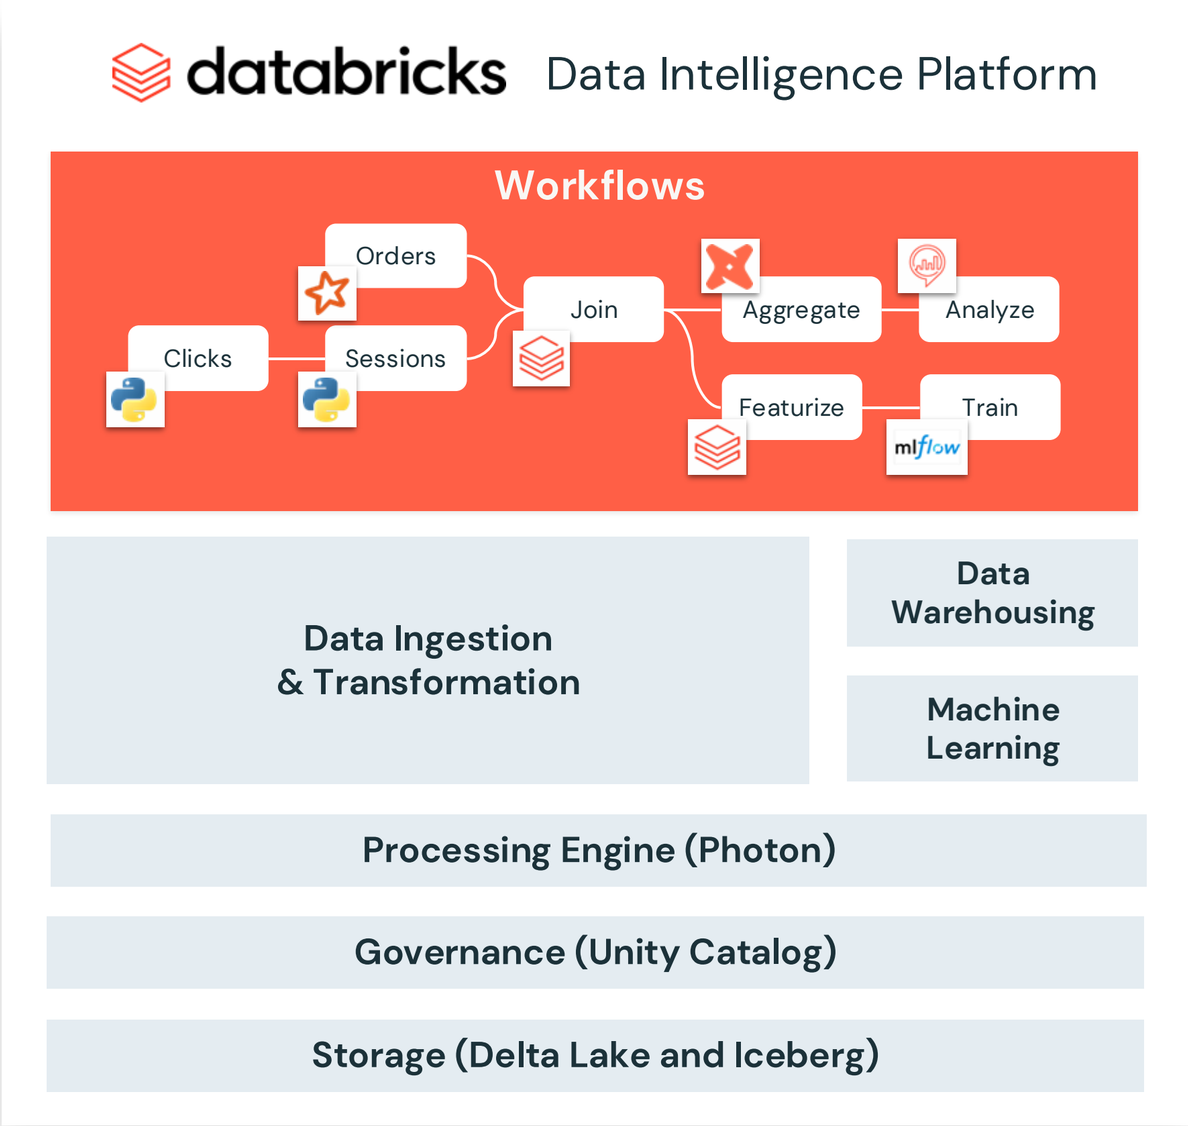

Jobs

Unified orchestration for data, analytics, and AI on the Data Intelligence Platform

- Simple authoring
- Actionable insights
- Proven reliability

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lakeflow Jobs

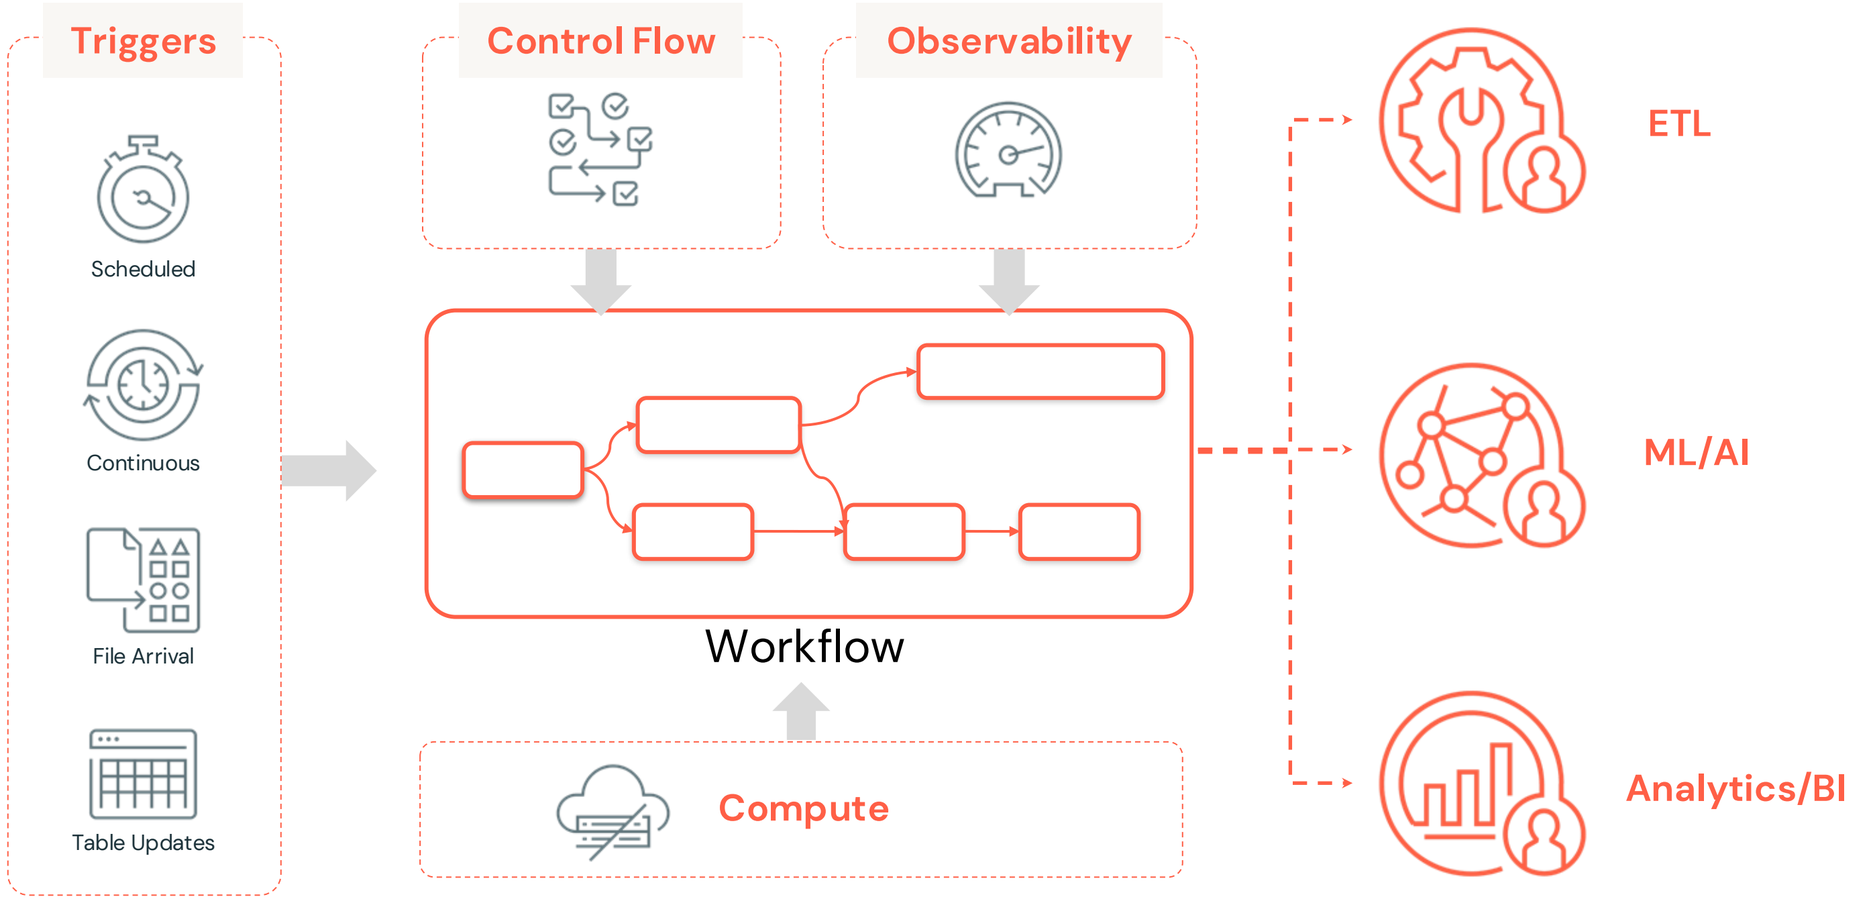

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lakeflow Jobs
Core Components

Deploy workload with lakeflow jobs

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Section Learning Objectives

- Understand the building blocks of Lakeflow Jobs, including the concepts of jobs and tasks.
- Identify and describe the various task types and configuration options available in Lakeflow Jobs.
- Explain how Directed Acyclic Graphs (DAGs) enable task orchestration and workflow patterns.
- Learn how to create and configure jobs using the Lakeflow Jobs UI.

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Agenda

Section Overview – Lakeflow Jobs Core Components
• Building Blocks of Lakeflow Jobs
• Task Orchestration
• Creating a Job Using the Lakeflow Jobs UI

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lakeflow Jobs Core Components

LECTURE

Building Blocks of Lakeflow Jobs

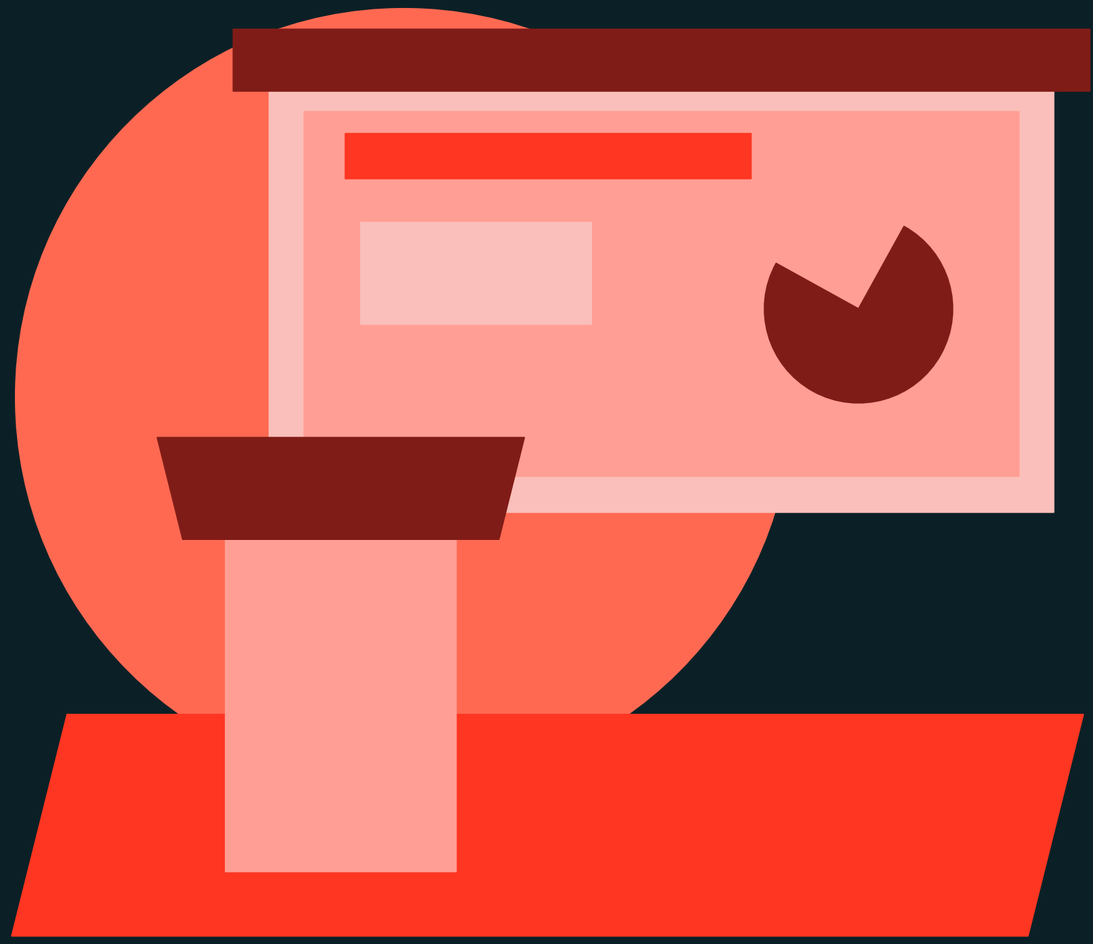

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Building Blocks of Lakeflow Jobs

Job and Tasks

- A **Job** is the primary resource for:
  - Scheduling
  - Coordinating
  - Running operations such as data processing, ETL, analytics, and machine learning workloads within the Databricks environment.

- A **task** is a **single unit of work** within a job that executes a specific workload such as a notebook, script, query and more.
- Each **job** consists of **one or more tasks**, which are the individual units of work that make up the job

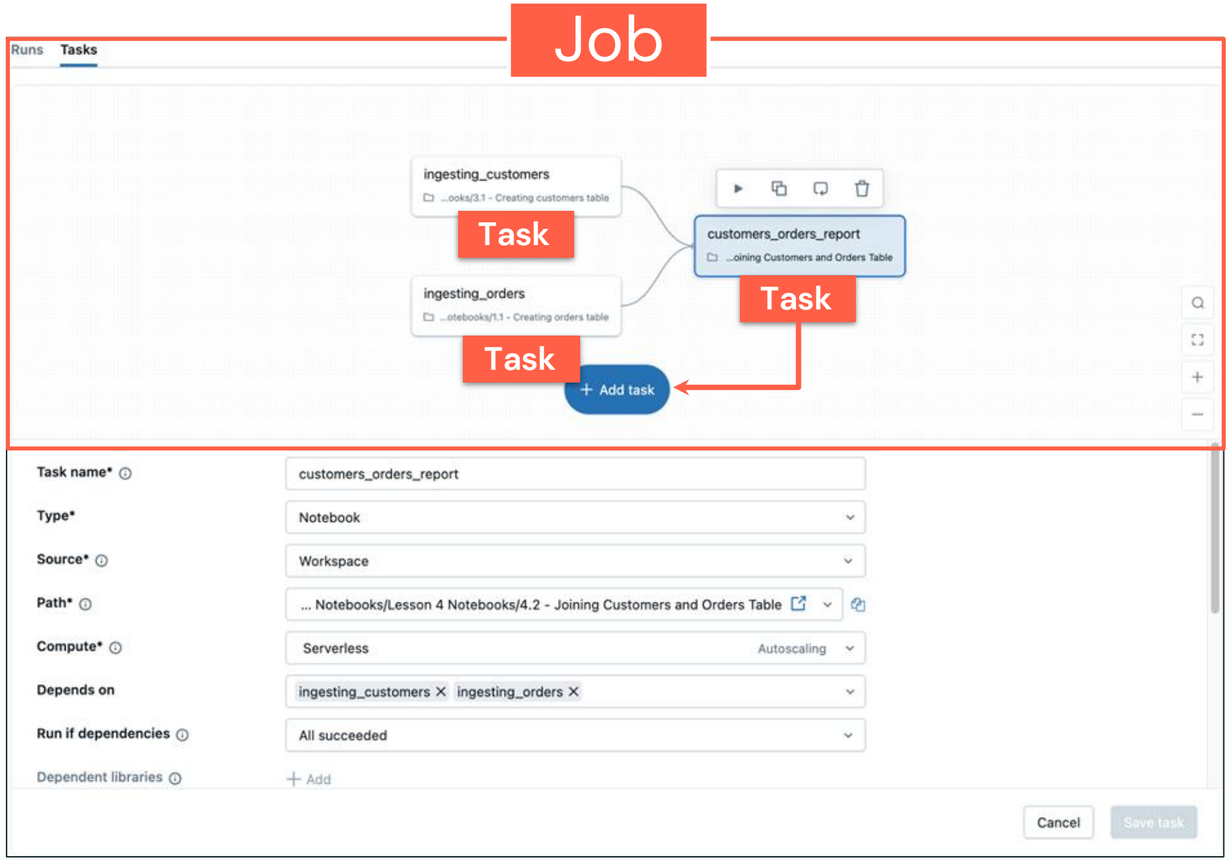

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Building Blocks of Lakeflow Jobs

Jobs consist of one or more **Tasks**

- Databricks Notebooks
- Spark Declarative Pipelines
- SQL Files/ Queries
- Python Scripts
- Python Wheels
- AI/BI Dashboards
- Spark Submit
- dbt
- Java JAR file
- Power BI

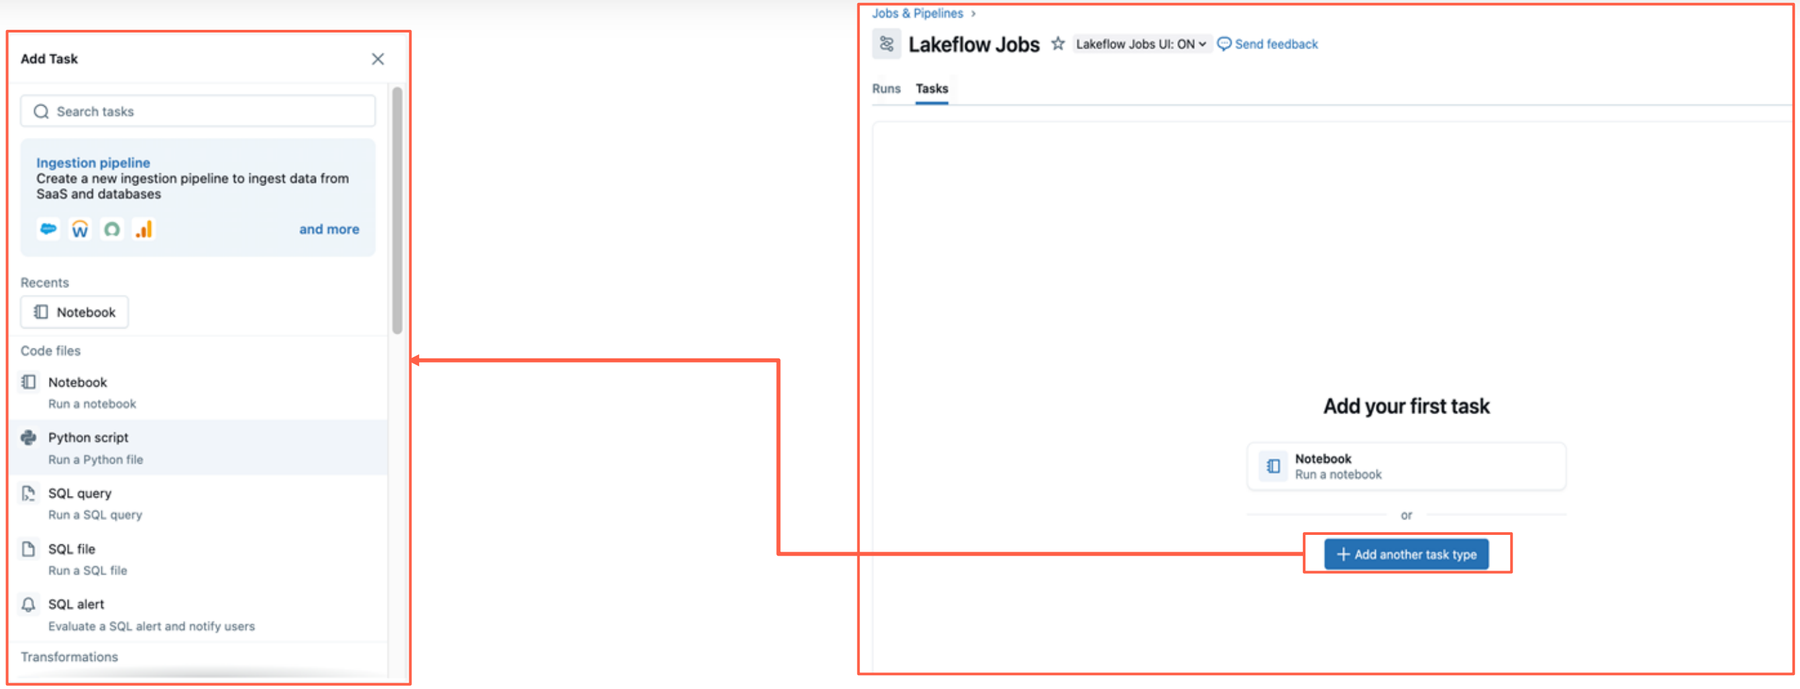

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Building Blocks of Lakeflow Jobs

Tasks Options

- With **Jobs**, we can set a **particular task**
- We have different **options available** based on the task:
  - Defining path
  - Adding libraries
  - Adding parameters
  - Enabling notification
  - All options are specific to the specific task

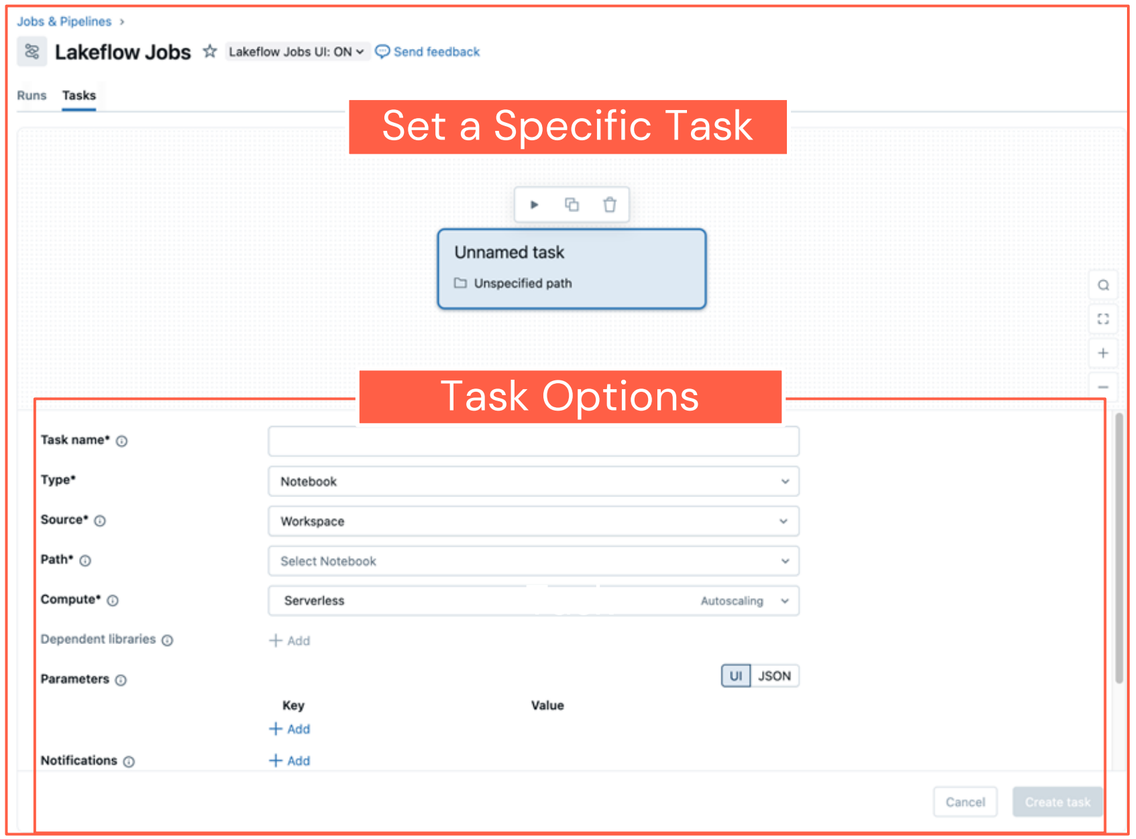

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Building Blocks of Lakeflow Jobs

Tasks – Notebook

Based on **task type selected,** you get specific options.

Example

For **Notebook** Tasks:

- Source
- Path
- Compute options
- And More

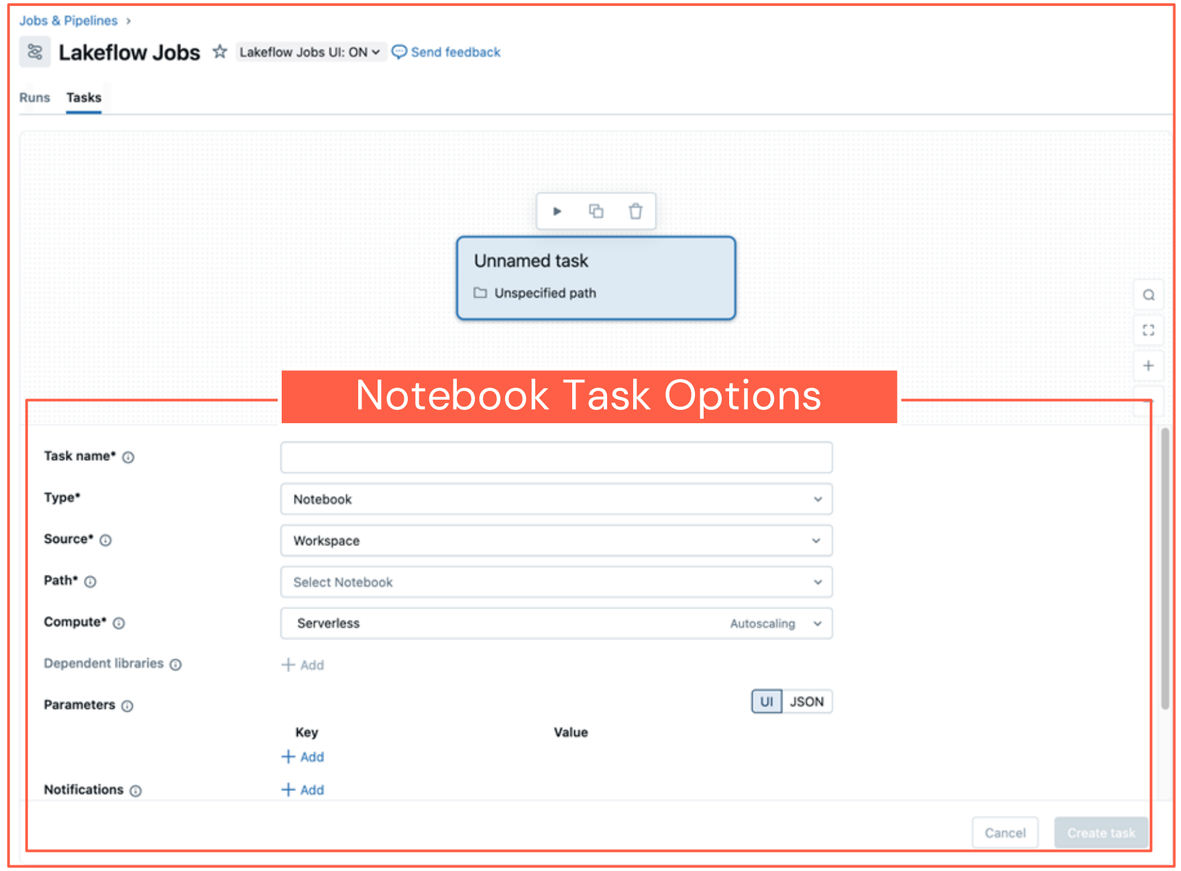

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Building Blocks of Lakeflow Jobs

Tasks – SQL

Based on **task type selected,** you get specific options

Example

For **SQL** Tasks:

- Task name
- SQL query
- SQL warehouse

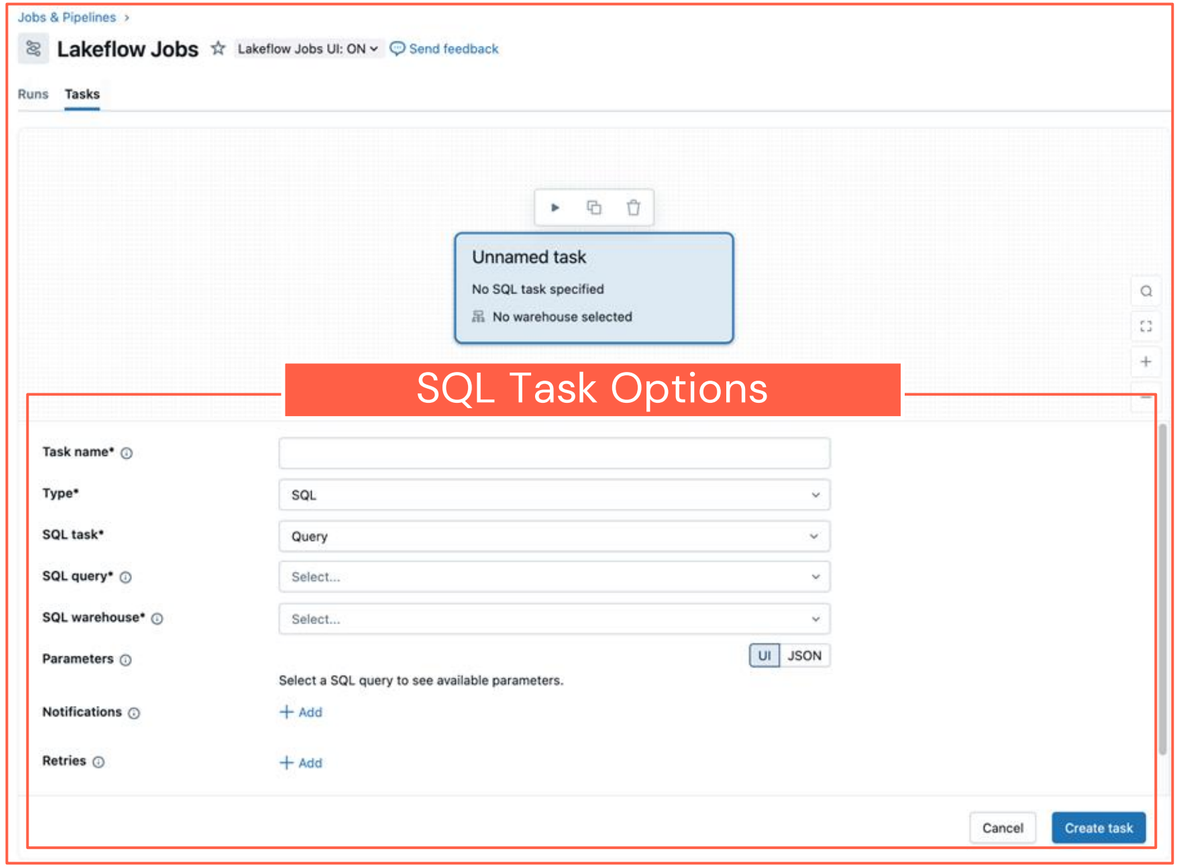

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Building Blocks of Lakeflow Jobs

Languages Supported

- Python
- SQL
- Scala
- R
- Java [via a JAR file]

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Building Blocks of Lakeflow Jobs

Jobs consist of one or more **Tasks**

- Databricks Notebooks
- DLT
- SQL Files/ Queries
- Python Scripts
- Python Wheels
- AI/BI Dashboards
- Spark Submit
- dbt
- Java JAR file
- Power BI

Control flows can be established between **Tasks**

- Sequential
- Parallel
- Conditionals (If/else)
- Run Job (Modular)
- For each

Jobs supports different **Triggers**

- Scheduled (Cron)
- API Trigger
- File Arrival Triggers
- Table Triggers
- Continuous (Streaming)
- Manual Trigger

COMPUTE

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Building Blocks of Lakeflow Jobs

Jobs can be executed on different compute

Development, Ad-Hoc Analysis & Exploration

Interactive Clusters

Interactive or all-purpose clusters can be shared by multiple users.

It is best for performing ad-hoc analysis, data exploration, or development.

Interactive should not be used in Production as they are not cost-efficient.

Production Grade & Operational Use-Cases

Job Clusters

Job clusters are approximately 50% cheaper as they terminate when the job ends, reducing resource usage and costs.

However, Job clusters are subject to cloud providers start-up times.

With Databricks Jobs, you can re-use the same cluster across Tasks for better price performance!

Simple, Fast, Reliable, Cost Efficient!

Serverless

Serverless Workflows are a fully managed service that are operationally simpler and more reliable

They provide you with faster clusters & auto-scaling capabilities providing you with a better user experience for a lower cost.

With out of the box performance optimizations, Serverless provides you with a lower overall TCO.

Compute for SQL Queries and BI, Serverless by default

SQL Warehouse

Purpose-built for SQL queries, dashboards, and BI; Can be attached to notebook as well.

High concurrency + autoscaling via Intelligent Workload Management for consistent low latency.

auto-start/auto-stop and adjustable cluster size and max clusters for peak loads helps control costs.

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Building Blocks of Lakeflow Jobs

Serverless Performance Mode

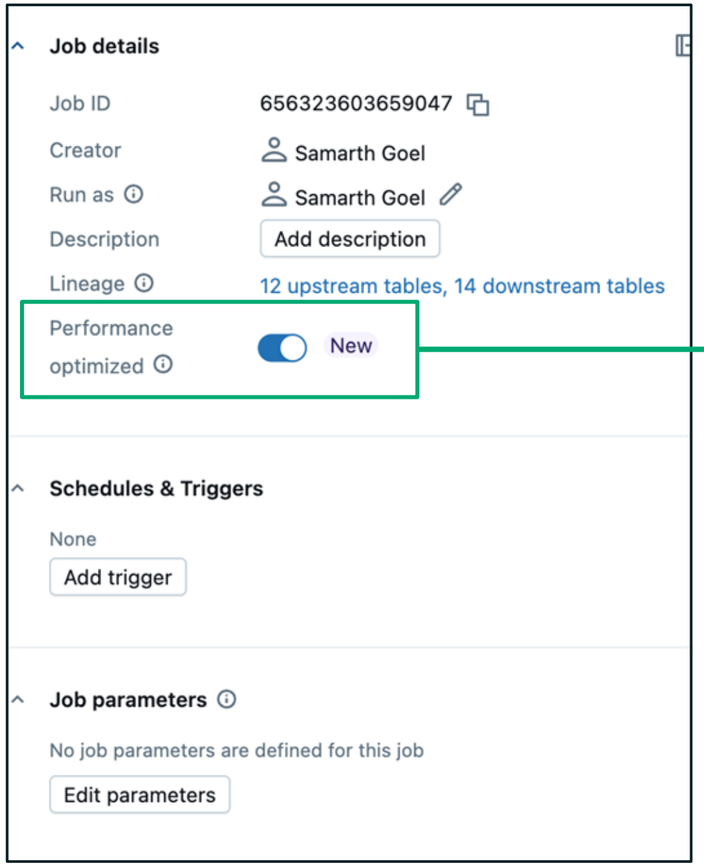

Use the **Performance optimized** setting in the job details page to choose between **lower cost** and **faster execution** for serverless tasks.

Performance Optimized (Off)

- Standard mode focuses on **cost-efficiency**
- **Longer startup** time (typically 4–6 minutes)
- Best for **non-urgent workloads** with flexible timing

Performance Optimized (On)

- **Enables faster** job startup and execution
- Ideal for **time-sensitive** workloads
- **Applies** only to tasks with **serverless** compute in job

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Building Blocks of Lakeflow Jobs

Selecting Compute

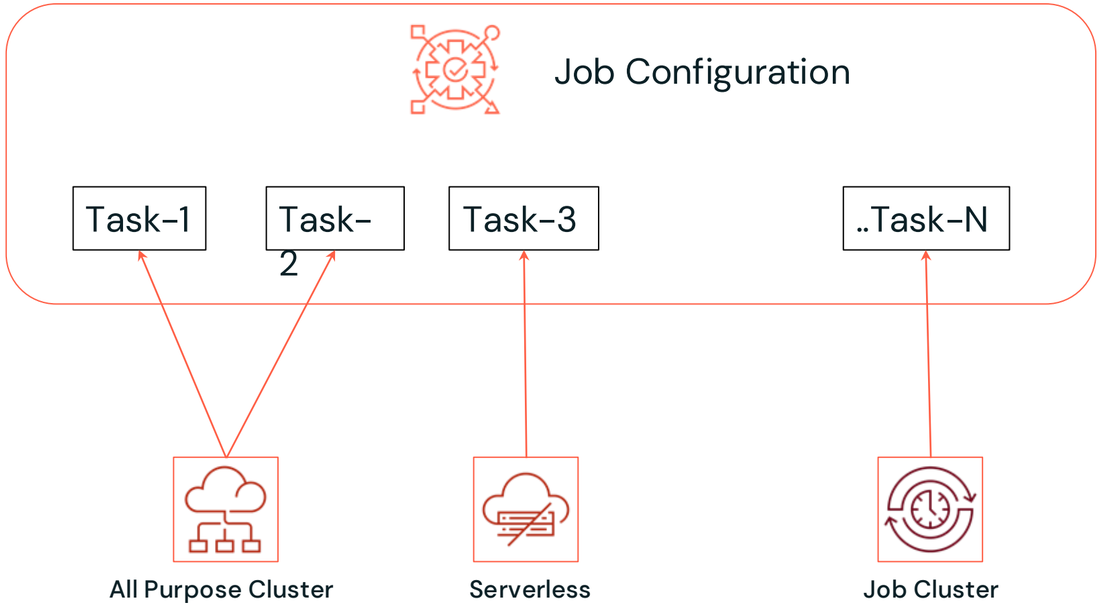

- A Job can have **one or more Tasks** inside it
- Each Task can be assigned its **own compute resource**
- Tasks in the same Job can either:
  - **share** the same compute or
  - use **different** compute as required

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lakeflow Jobs Core Components

LECTURE

Task Orchestration

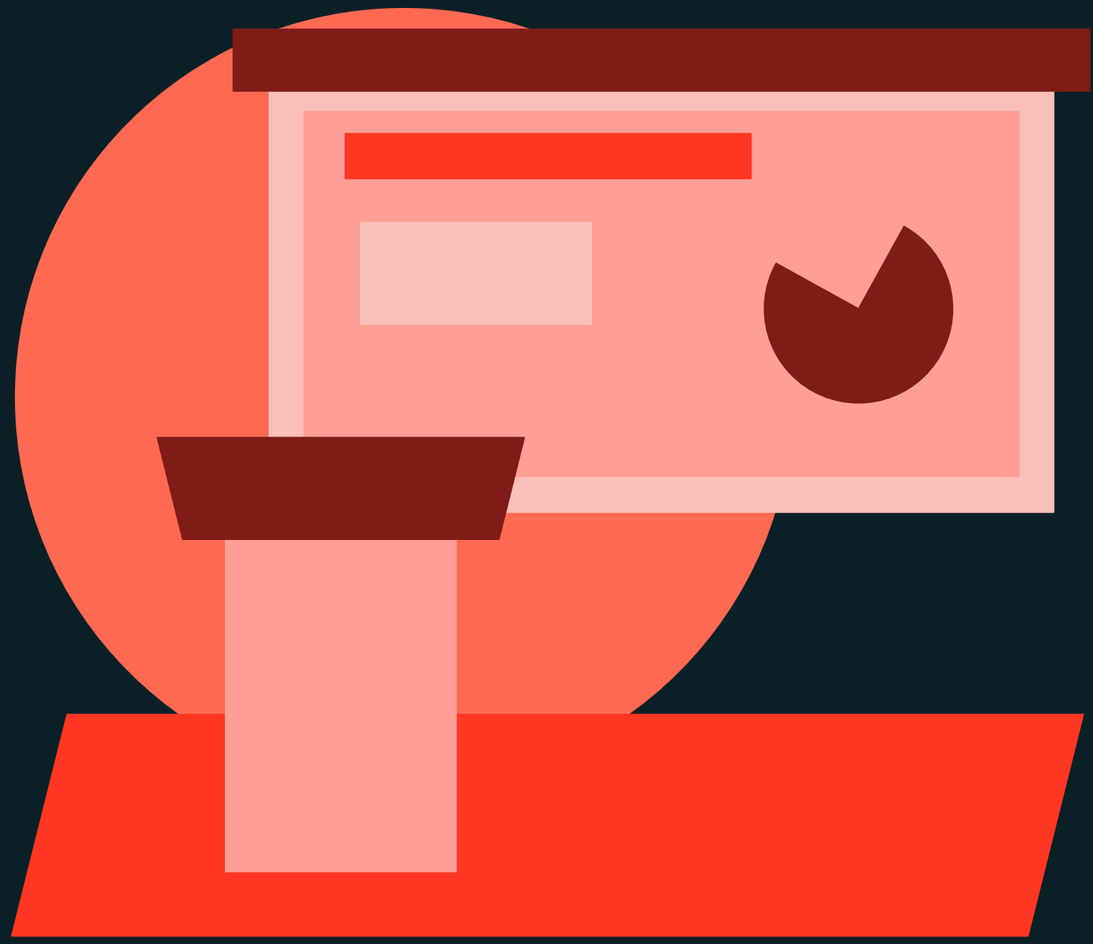

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Task Orchestration

What is a DAG?

A DAG is a **conceptual representation** of series of activities, including data processing flows.

- **D**irected - unambiguous direction for each edge
- **A**cyclic - contains no cycles
- **G**raph- collection of vertices connected by edges.

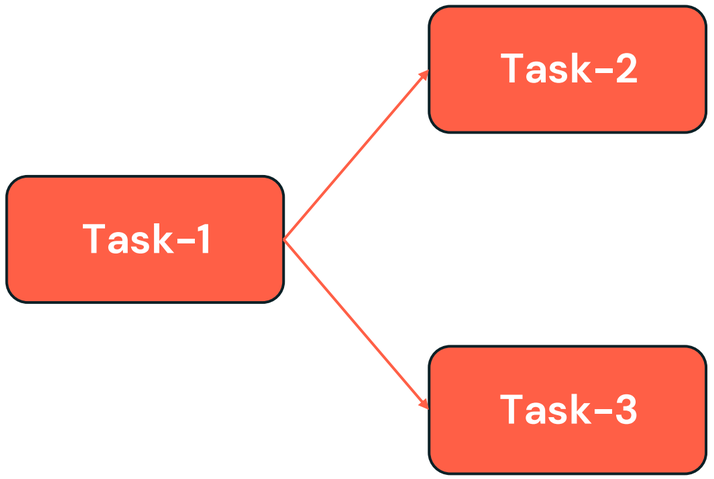

Diagram of a DAG = vertices
edges =

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Task Orchestration

Task Orchestration Overview

- Databricks Jobs supports **task orchestration**
  - The ability to run multiple tasks as a Directed Acyclic Graph (DAG)
  - Orchestrate tasks using the **Databricks UI** , **API, SDK** or **Databricks Asset Bundles**
- You can define the **order of execution of tasks** in a job by configuring task dependencies, creating a **DAG of task execution**

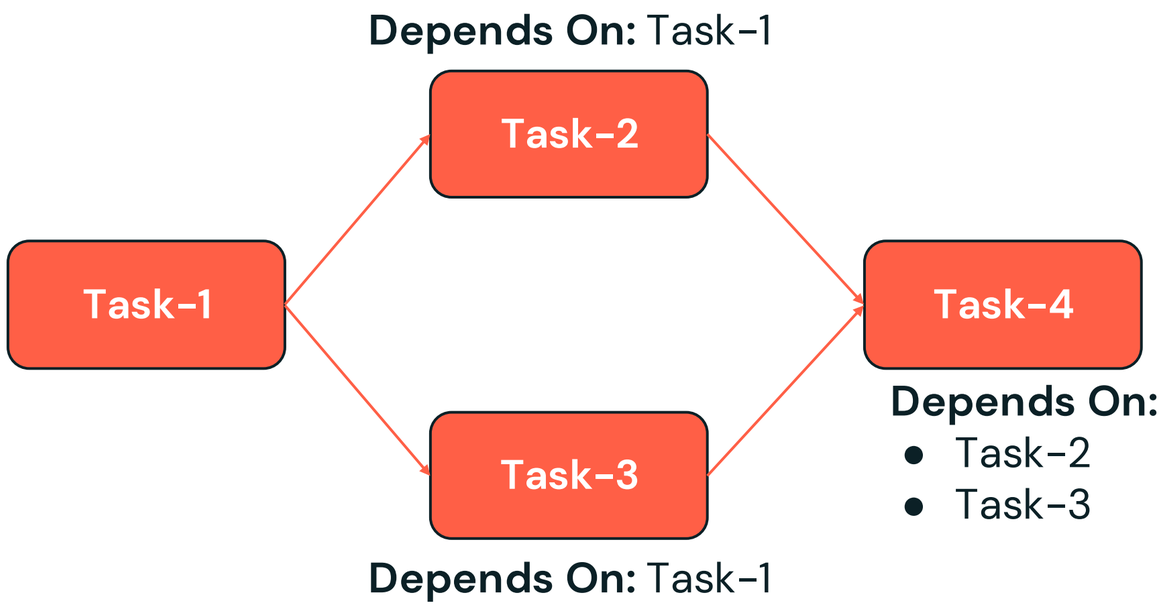

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Task Orchestration

Common Workloads Patterns

Sequence

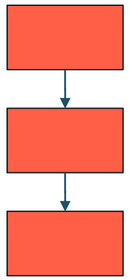

Sequence

- Data transformation/ processing/cleaning
- Bronze/silver/gold tables

Funnel

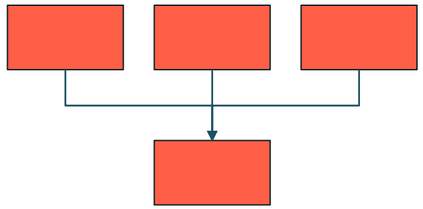

Funnel

- Multiple data sources
- Data collection

Fan-out

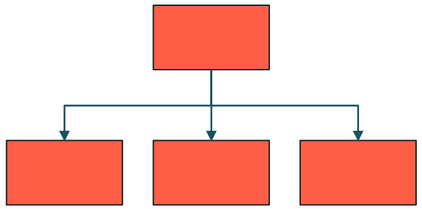

Fan-out, star pattern

- Single data source
- Data ingestion and distribution

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Introduction to Data Engineering in Databricks

LECTURE

Course Project Overview

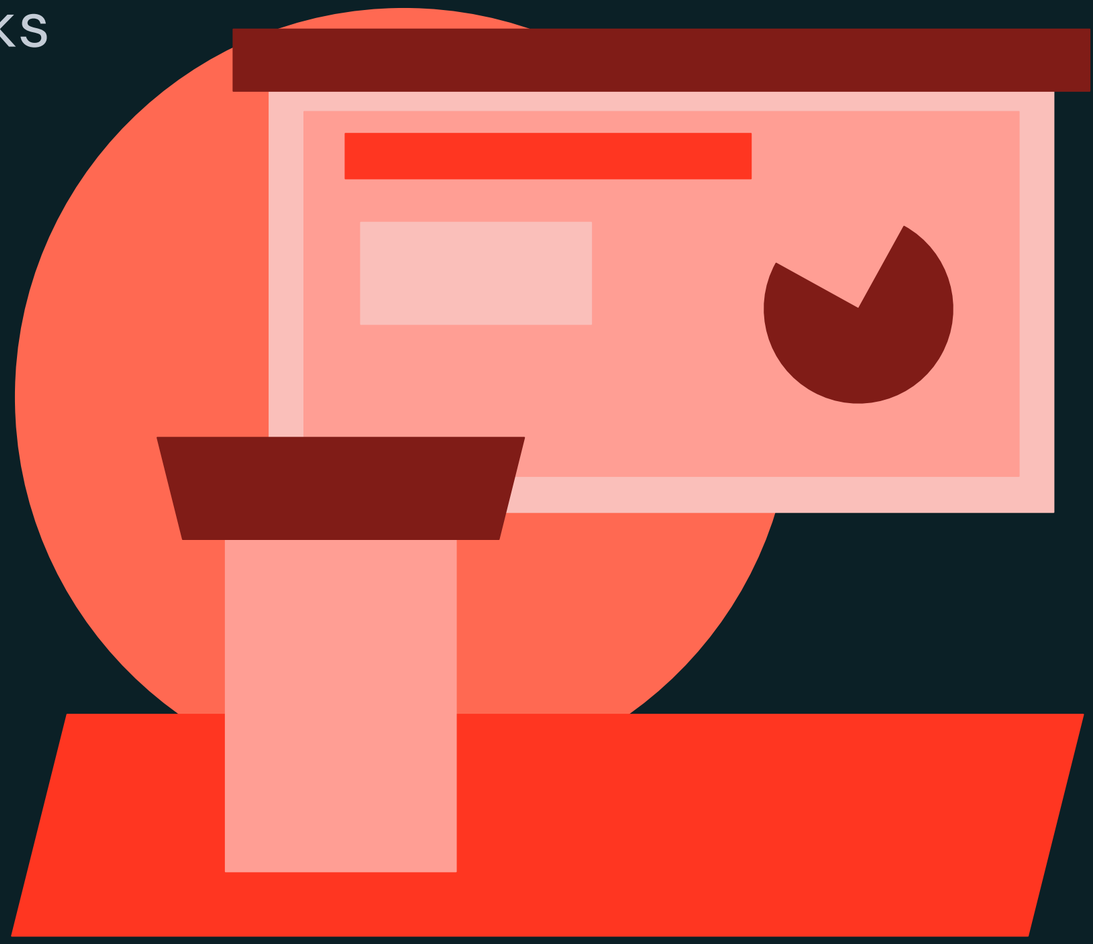

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Course Project Overview

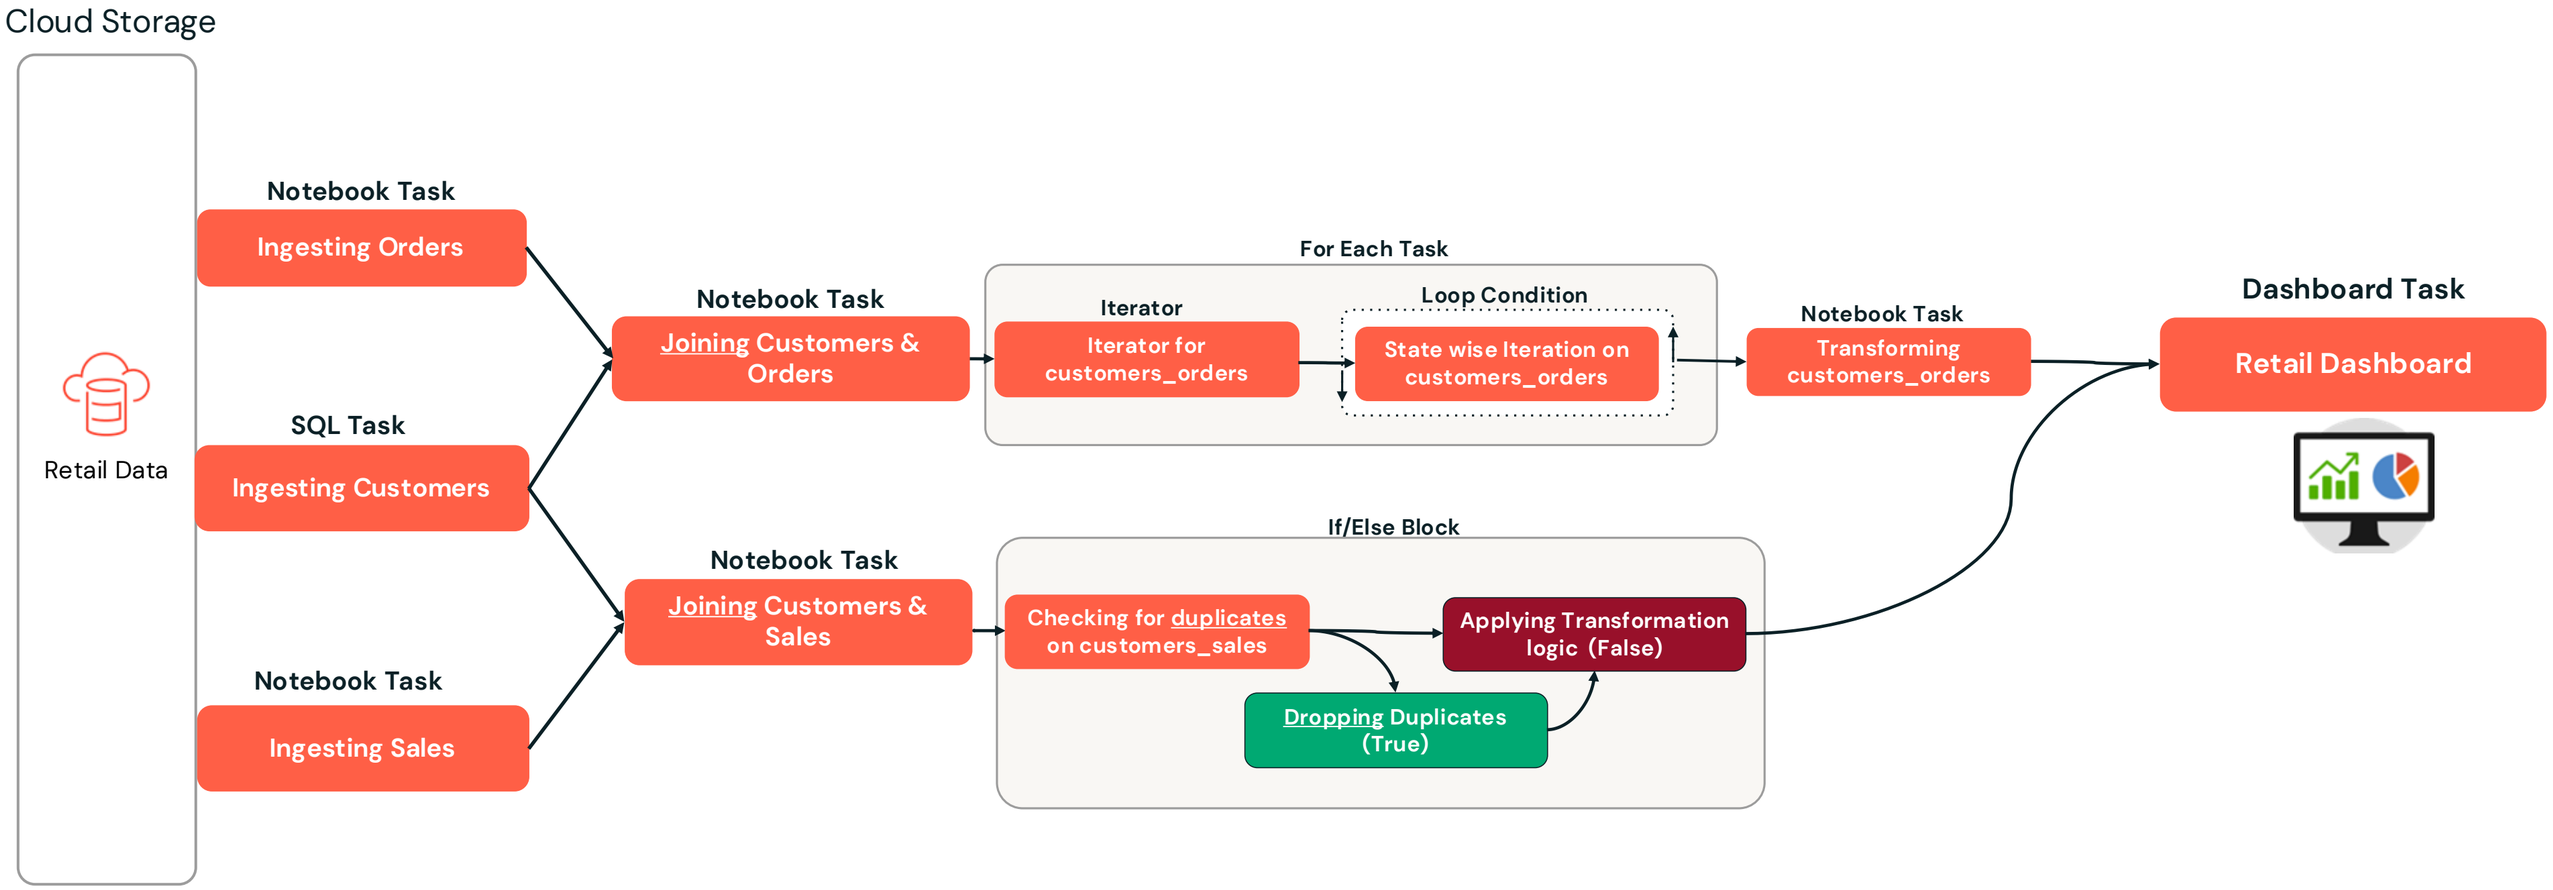

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lakeflow Jobs Core Components

DEMONSTRATION

Creating a Job Using the Lakeflow Jobs UI

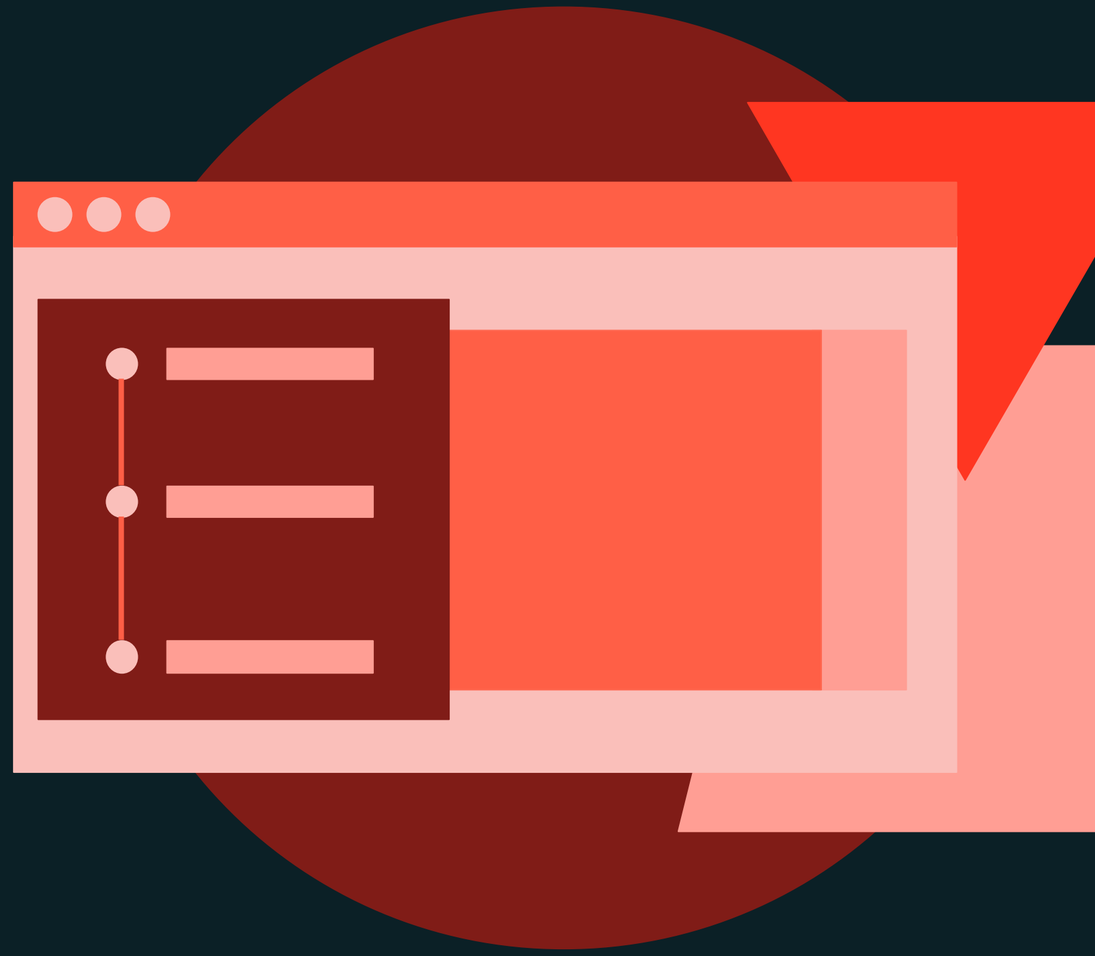

1 Demo - Creating a Job Using the Lakeflow Jobs UI

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lakeflow Jobs Core Components

LAB EXERCISE

Create your First Job

2 Lab - Create your First Job

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Creating and Scheduling Jobs

Deploy Workloads with Lakeflow Jobs

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Section Learning Objectives

- Configure task parameters,task values,dynamic value references, notifications, and retry policies for Lakeflow Jobs.
- Set up and manage different job scheduling options, including scheduled, file arrival, Table update and continuous triggers.
- Automate job execution using triggers to support time-based and event-driven workflows.
- Explore and apply scheduling options in the Databricks Lakeflow Jobs UI through hands-on demonstration.

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Agenda

Creating and Scheduling Jobs
• Common Task Configuration Options
• Jobs Schedulers and Triggers
• Demo – Exploring Scheduling Options

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Creating and Scheduling Jobs

LECTURE

Common Task Configuration Options

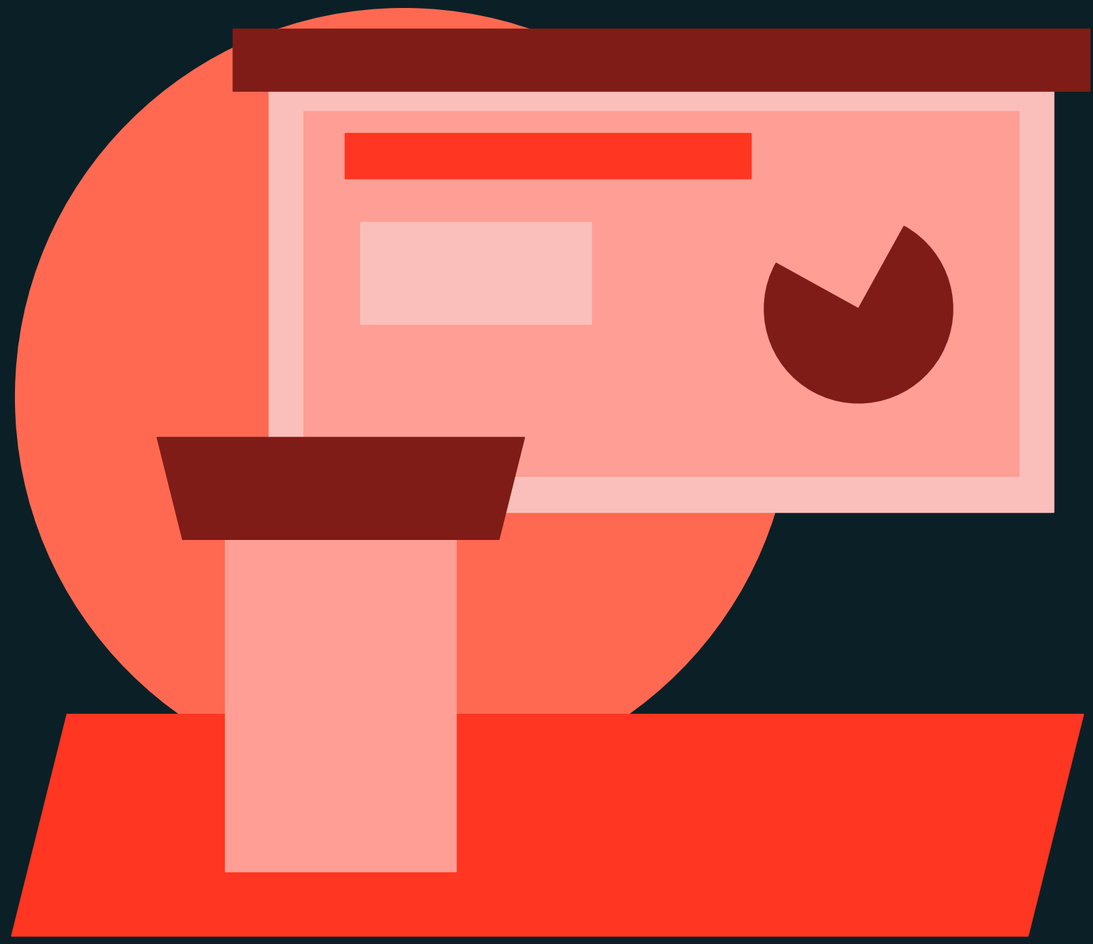

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Common Task Configuration Options

Overview

Parameters & Dynamic Value References

Notification Alerts

Retries

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Common Task Configuration Options

Parameter Overview

Task Parameters

A key-value pair defined at the **task level**

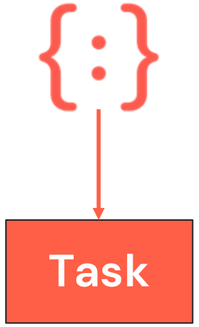

Job Parameters

A key-value pair defined at the **job level** and pushed down to all tasks

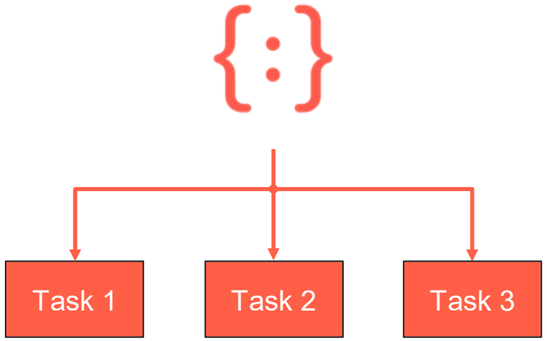

Job Parameters **override** Task Parameters when same **key** exists

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Common Task Configuration Options

Task Parameter Overview

- **Task Parameters** are **key-value** pairs that lets you pass values into tasks
- Support **advanced orchestration** such as
  - Conditional execution
  - Looping
  - Passing context between tasks

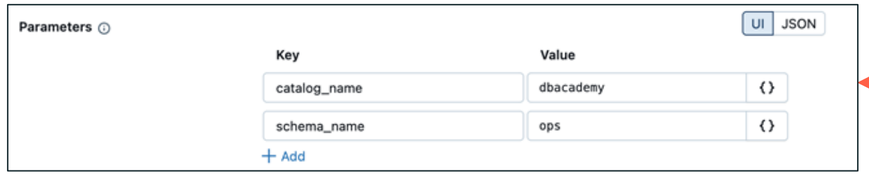

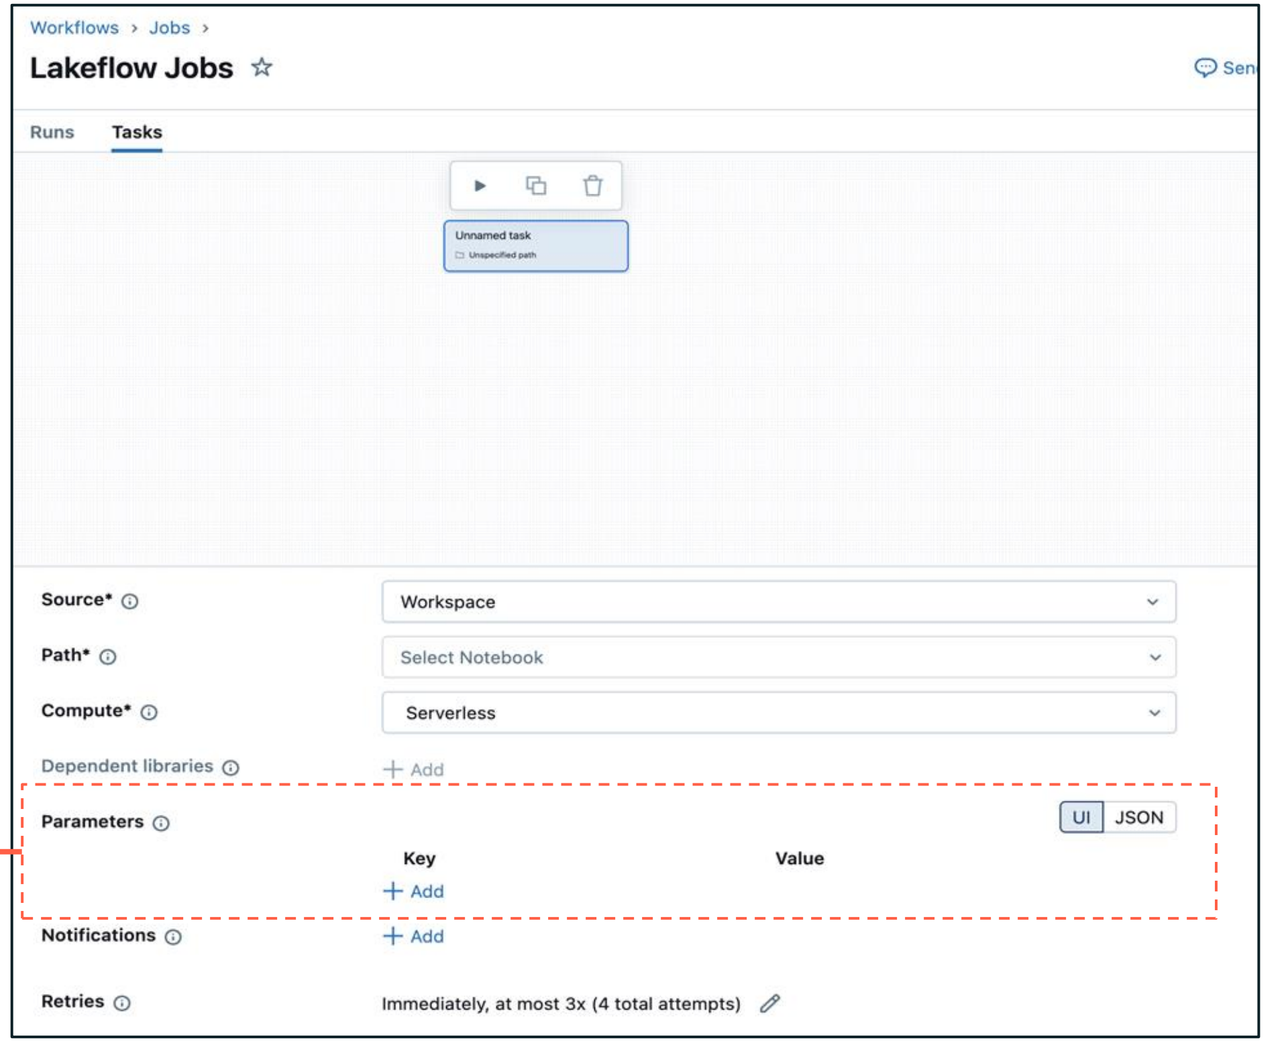

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Common Task Configuration Options

Job Parameter Overview

Job Parameters are **key-value pairs** defined at the **job level** that provide default values for the entire workflow.

- Applied to **all tasks** in the job automatically
- **Override task parameters** with the same key name
- Can be **overridden at runtime** when triggering job runs

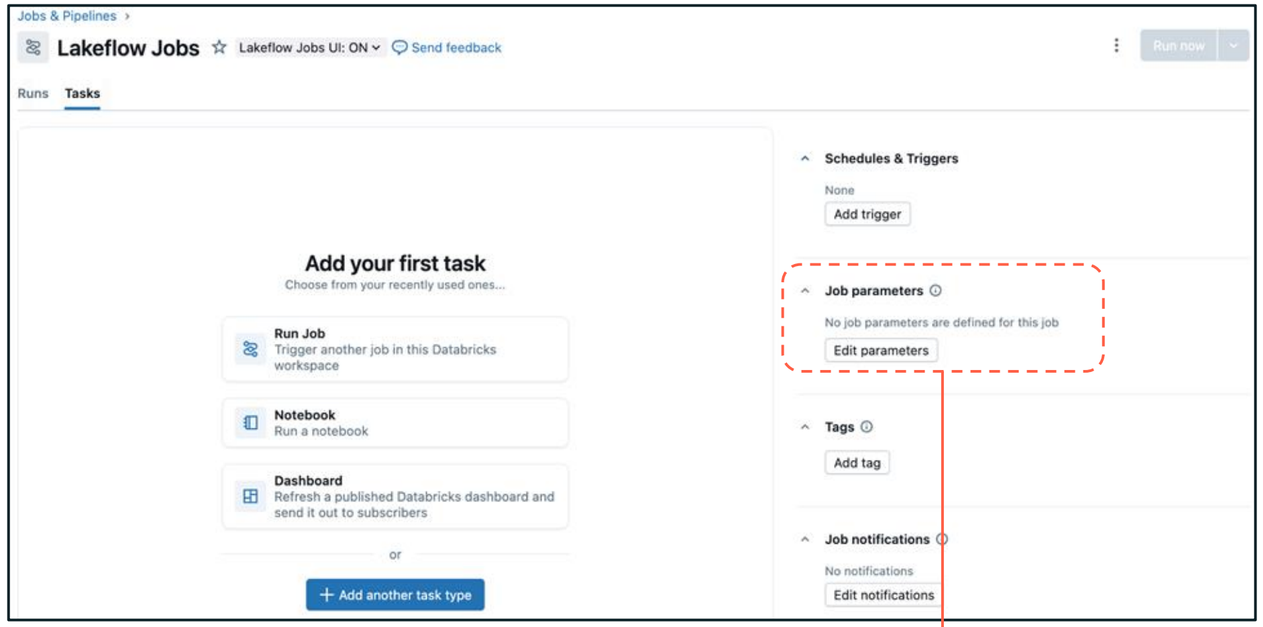

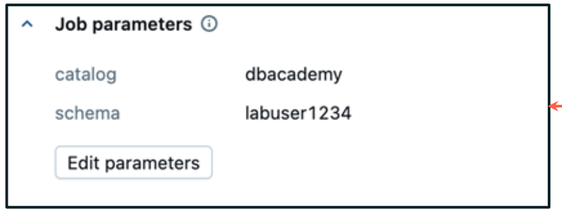

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Common Task Configuration Options

Setting and Accessing Parameters (UI)

Setting Parameter Values (UI)

Accessing Parameter Values in a Task

```
catalog = dbutils.widgets.get("catalog_name")

schema = dbutils.widgets.get("schema_name")
```

Setting Job Parameters:

- In your **job**, go to the **Parameters section**, and add key–value pairs at the job level.

Setting Task Parameters:

- Go to the **specific task in the job**,
- Go to **parameters** within the **task configuration** (alongside task name, type, and path)
- Add the key–value pairs

- Access both **task and job parameters** in notebook tasks with **`dbutils.widgets.get`**
- Retrieval methods vary by **task type** (e.g., notebook vs. SQL vs. Python wheel)

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Common Task Configuration Options

Dynamically Setting and Accessing Task Parameters

Dynamically Set Task Values with Code

```
dbutils.jobs.taskValues.set(
                key = "catalog_name",
                value = "dbacademy")
```

Task Values are **dynamic key-value pairs** that tasks create and share during **workflow execution**.

Accessing a Parameter from another Task

```
dbutils.jobs.taskValues.get(
            taskKey="task-name",
                  key="catalog_name")
```

You Reference specific upstream task parameter values into **downstream tasks**

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Common Task Configuration Options

Dynamic Value References

- **Dynamic Value References** allow referencing values at **runtime** from the job and task context (e.g., task outputs, parameters, retask values, registered task values, run metadata).
- These references use `{{ }}` notation and allow workflows to adapt to different execution.
- Common examples:
  - `{{job.start_time.day}}` - For accessing the day
  - `{{task.name}}` - For Accessing Task Name
  - `{{tasks.<task-name>.values.my_value}}`

    For accessing output from an upstream task

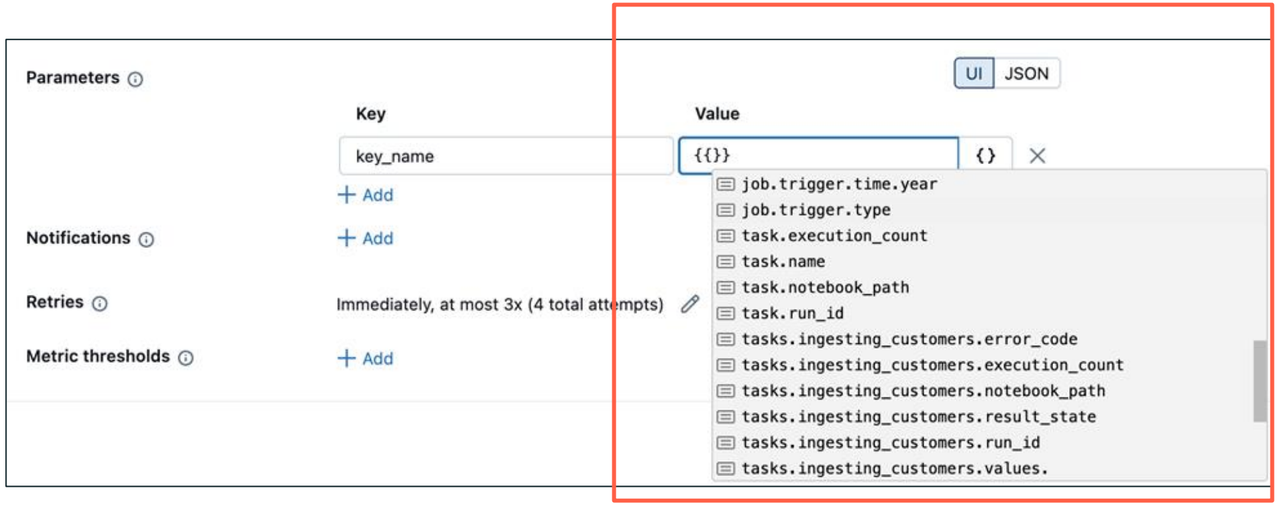

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Common Task Configuration Options

Notification Alerts

Notification Configurations

Job Level Notifications

- A notification alert is sent after **successful completion of the job**
- This setting can be adjusted in the Job Details section (right-side pane).

Task Level Notifications

- A notification alert will be sent when the **task is successfully completed**
- This notification setting can be customized at the task level.

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Common Task Configuration Options

Notification Alerts

- **Multiple** supporting destinations as Emails, Teams, pagerDuty,Slack and Webhook
- **Each task in a Job** could be configured **differently** to send notification
- Jobs triggers **notifications** when the task begins, completes or fails
- These notification could also be triggered in the case of **late Jobs:**
  - For duration threshold warning/timeout alerts
  - Streaming backlog(Issue with getting streaming data)
- Webhook allows custom integration with third-party API services

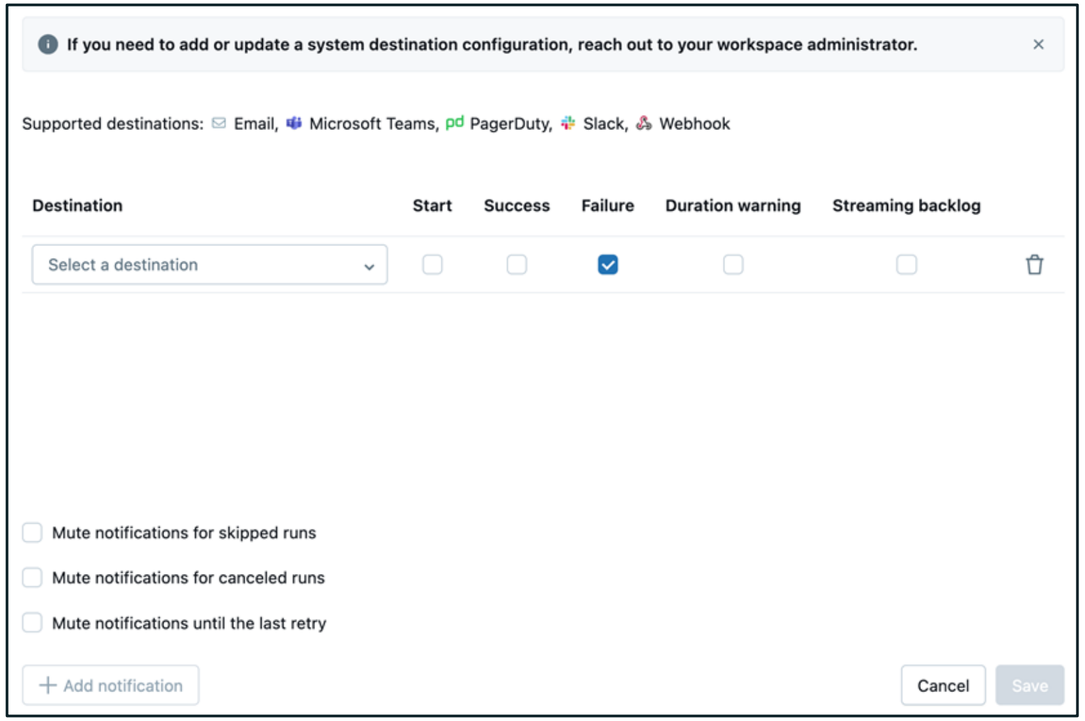

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Common Task Configuration Options

Retry Policy

A policy that determines when and how many times failed runs are **retried**

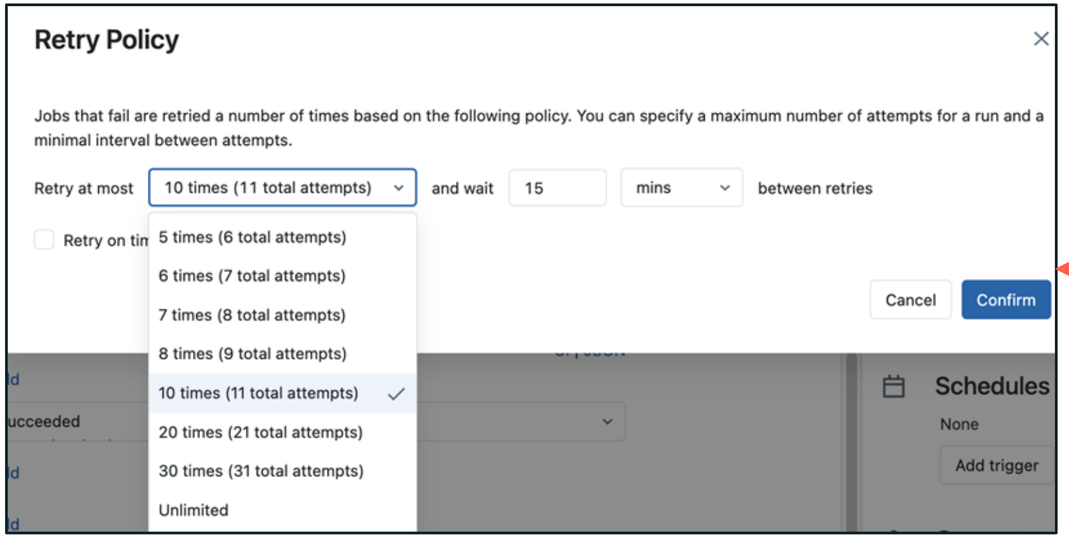

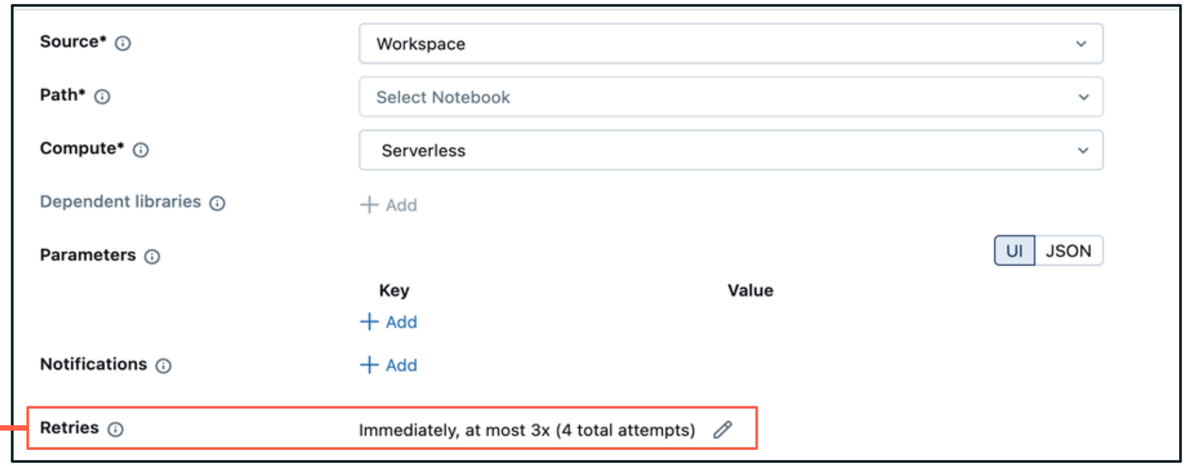

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Creating and Scheduling Jobs

LECTURE

Job Schedules and Triggers

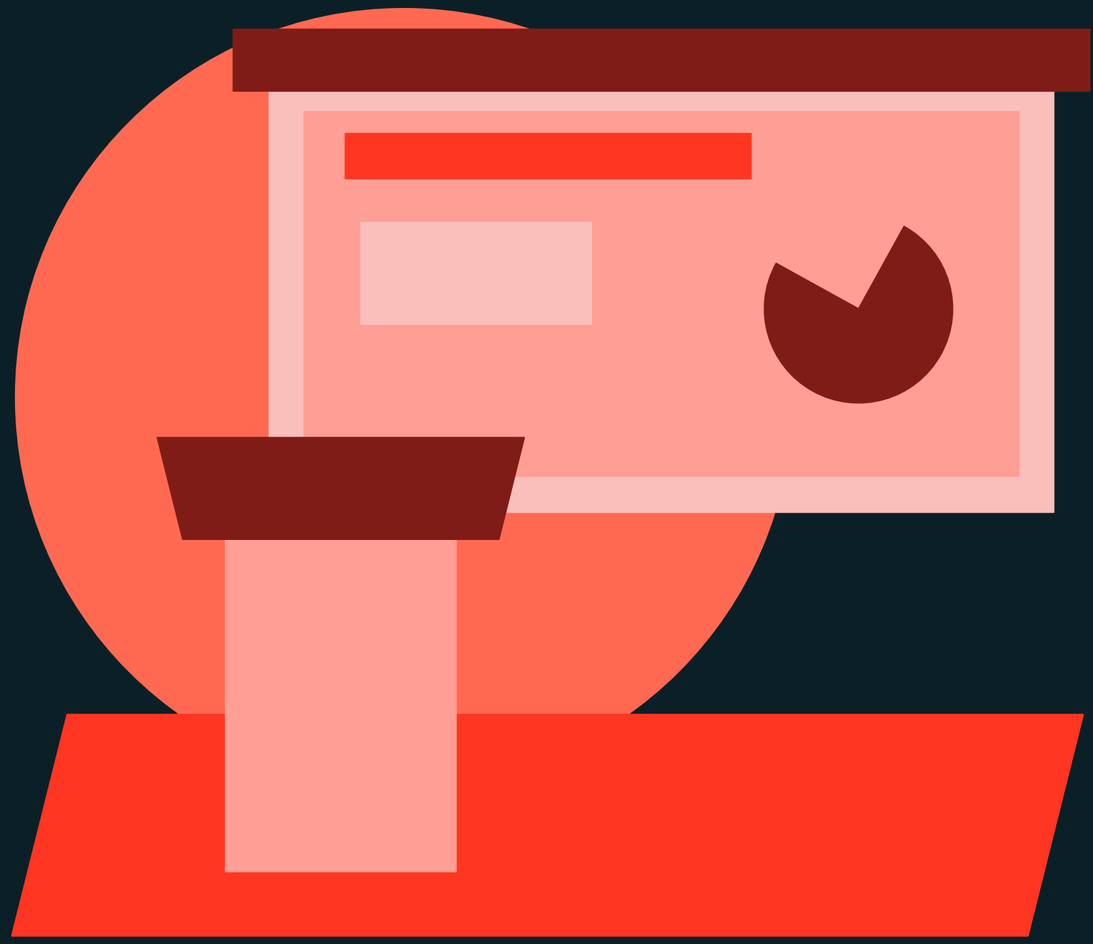

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Job Schedules and Triggers

Overview

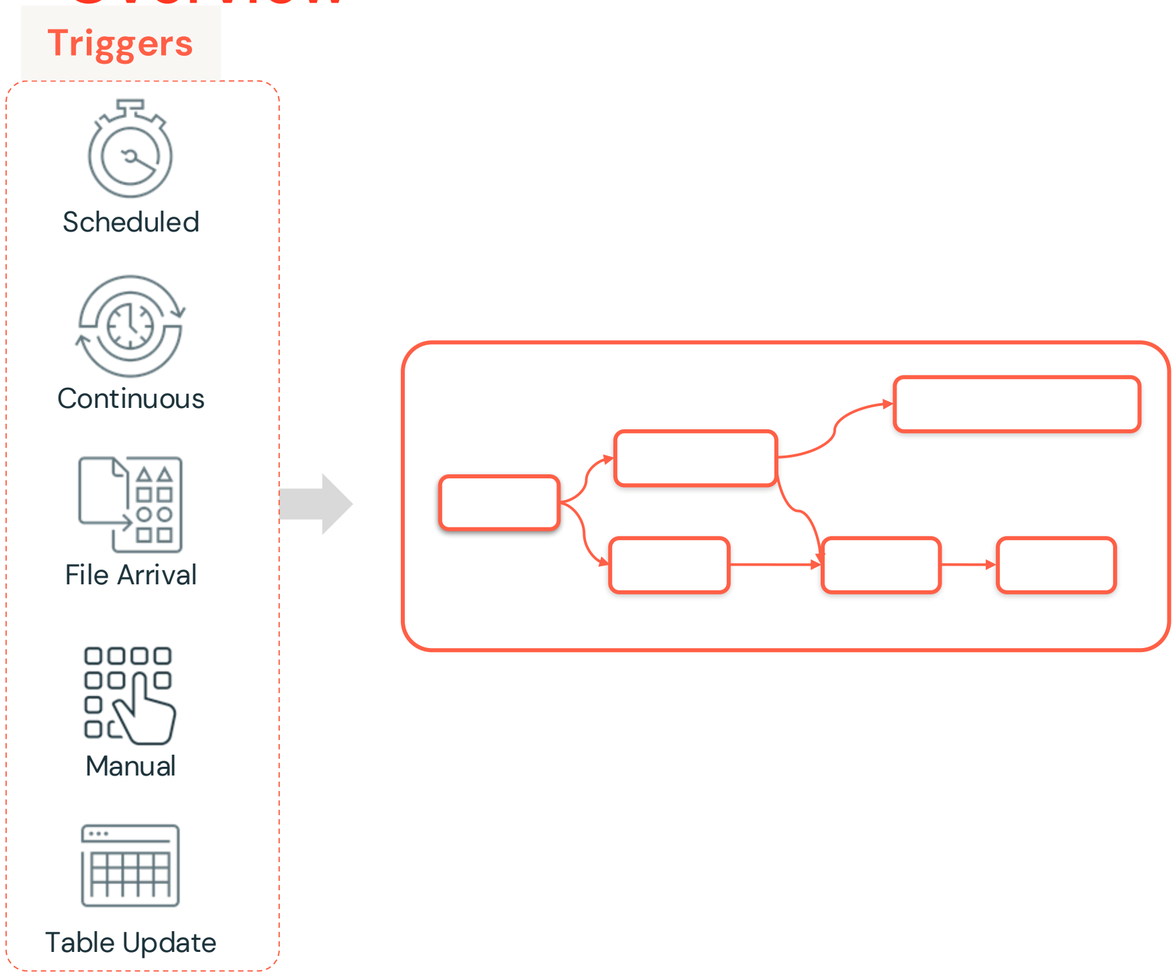

- A **trigger** is a rule that automatically starts a job run based on a **specific condition** or **schedule**
- **Common** trigger types include:
  1. Time-based schedules
  2. Continuous (always-on) execution
  3. File arrival events
  4. Manual trigger
  5. Table Update
- Triggers **enable automation**, so jobs can run without manual intervention

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Job Schedules and Triggers

Types of Triggers

1. Scheduled Trigger

- **Automatically run jobs** at set times or intervals, such as hourly, daily using the UI
- You can schedule jobs using **cron expressions** for automated execution
- They help **automate** recurring jobs, ensuring jobs run consistently without manual intervention

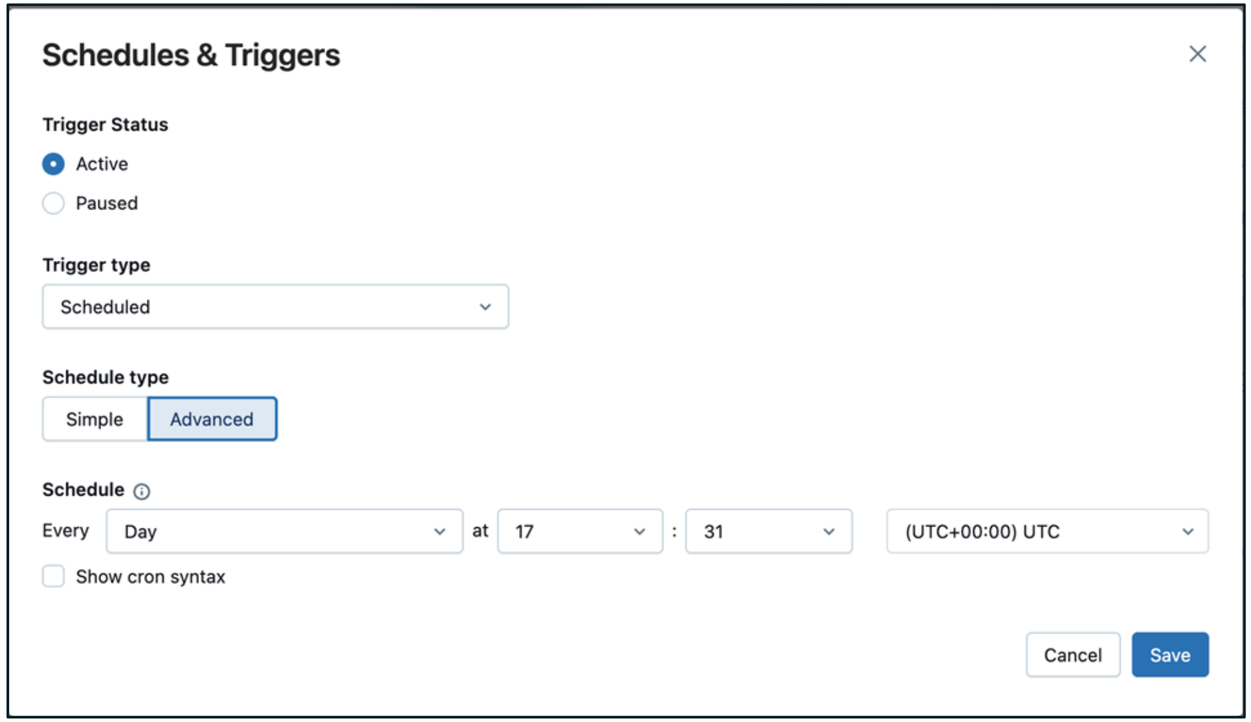

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Job Schedules and Triggers

Types of Triggers

2. File Arrival Trigger

- **Automatically** trigger jobs **when new files are detected** in a specified storage location, with support for:
  - AWS S3, Azure Storage, GCP GS and Databricks Volumes
- Enable **event-driven processing** to start data workflows upon **file arrival**
- Ideal for **automating jobs** with **unpredictable or irregular** data ingestion patterns

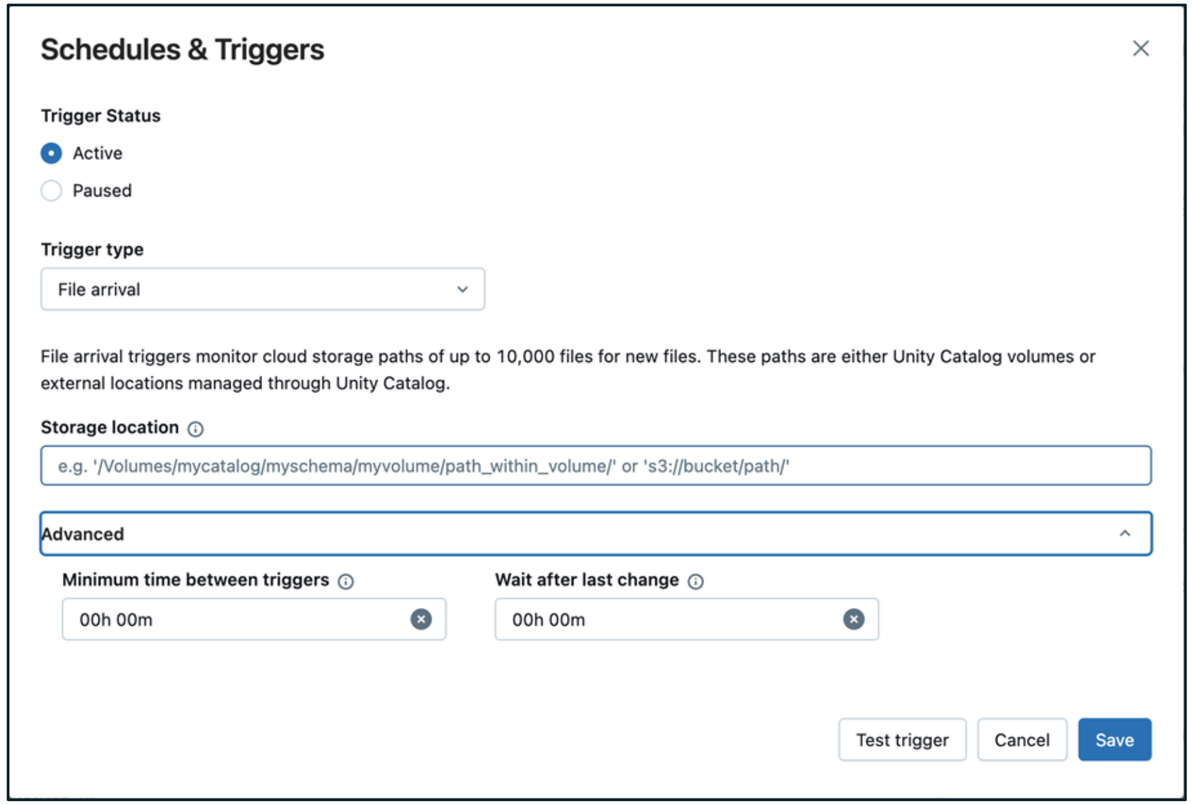

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Job Schedules and Triggers

Types of Triggers

3. Continuous Trigger

- **Continuously** runs jobs by starting a new run as soon as the previous one finishes or fails
- **Built-in retry** logic is automatically managed by Databricks
- Ideal for **streaming** workloads

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Job Schedules and Triggers

Types of Triggers

4. Manual Triggers

- **Manual (None) trigger** lets jobs run on demand without any schedule or event.
- Can be **started from UI** (Run now or Run now with different settings), API, CLI, SDK, or via DABS
- Best for **ad hoc runs, debugging, and one-off backfills**
- Can be combined with other trigger types later

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Job Schedules and Triggers

Types of Triggers

5. Table Update Trigger

- Table Update Trigger automatically starts a job when **specified tables** are updated .
- **Monitors** one or more tables for changes (insert, update, delete, or merge) .
- Helps **replace** time-based scheduling (like cron jobs) with real-time orchestration, ensuring jobs run as soon as **new data arrives .**

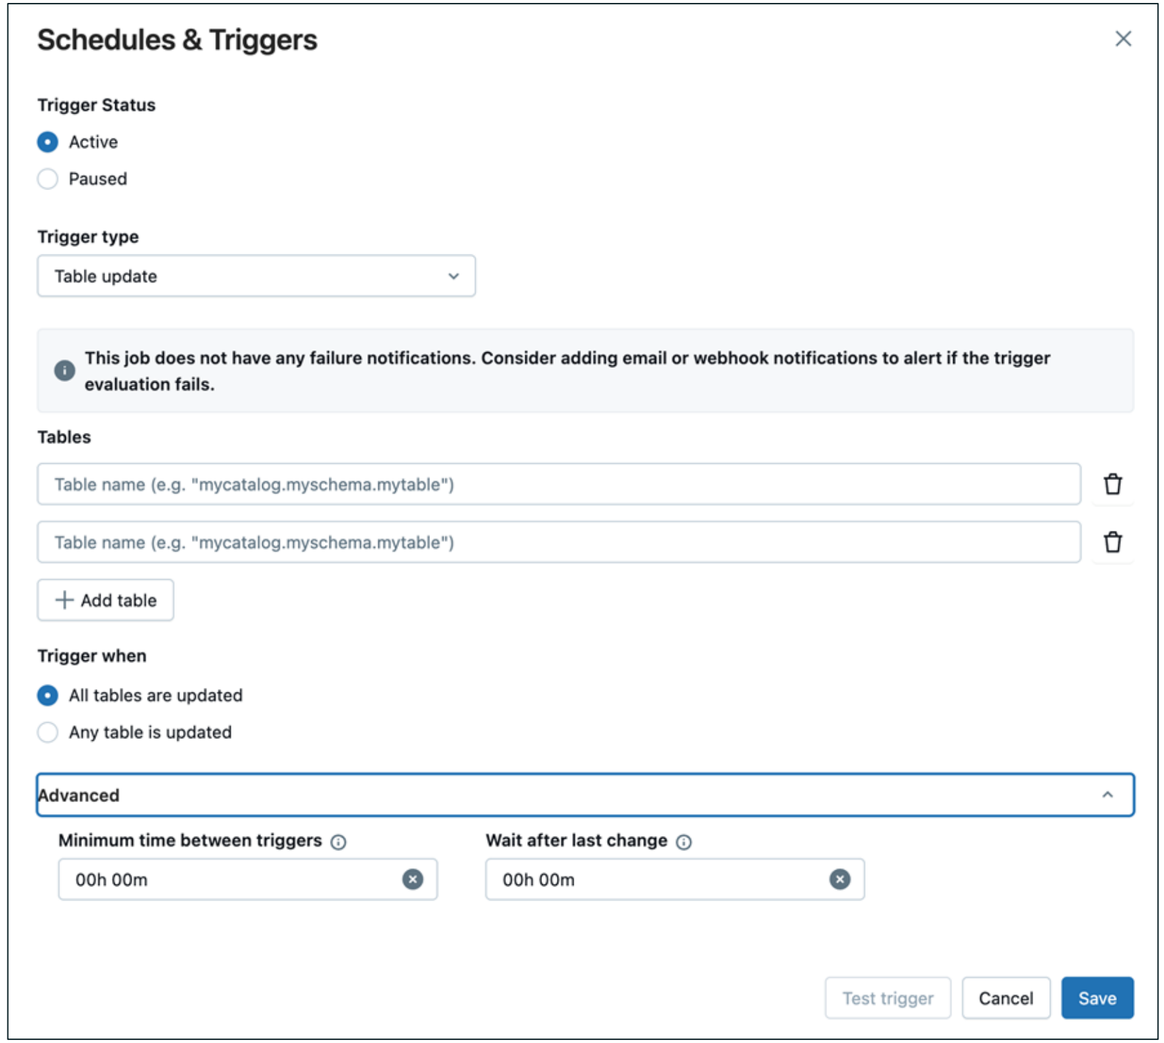

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Job Schedules and Triggers

Types of Triggers - Table Update

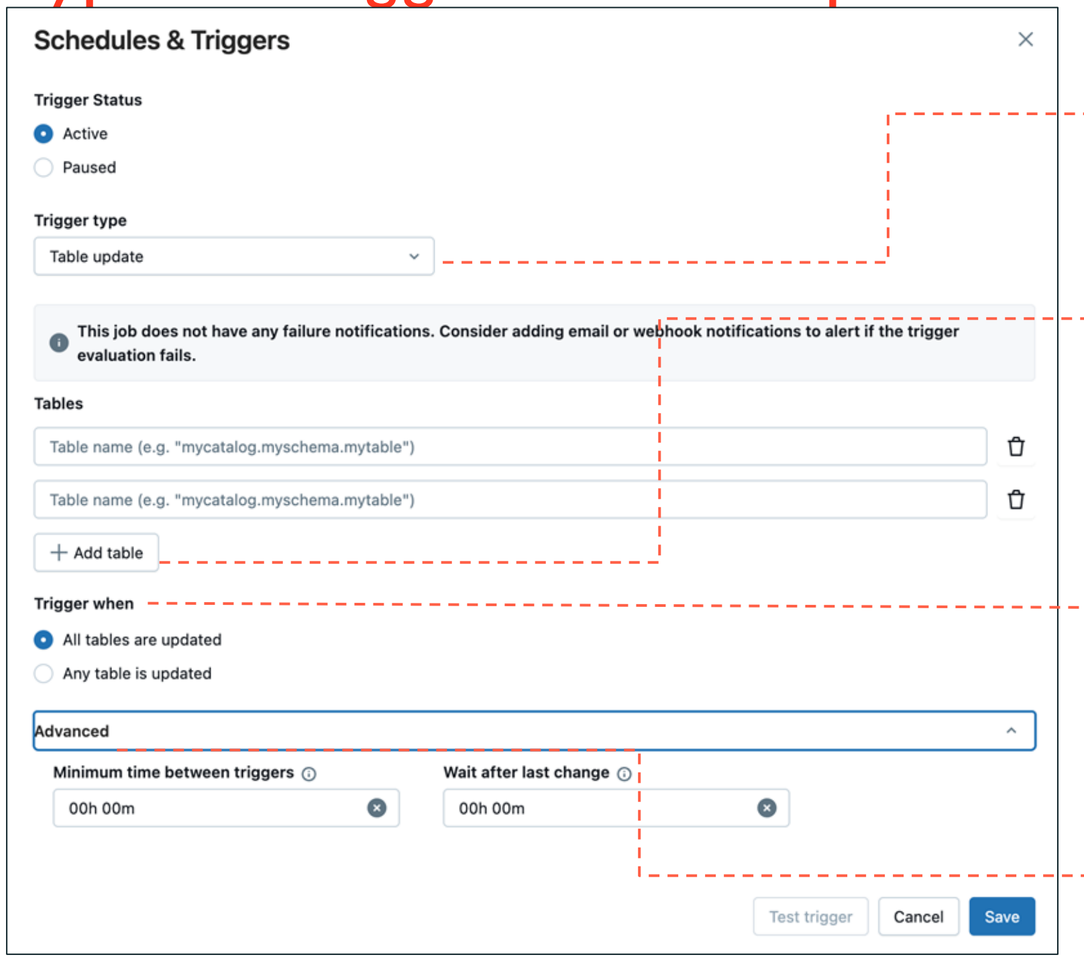

Step-1

Selecting trigger type

Step-2

Adding source tables

- You can select up to 10 tables per trigger.
- Works with Unity Catalog-managed Delta, Iceberg, external Delta tables, materialized views, and streaming tables.

Step-3

Defining triggering condition

- Any table updated – triggers when the first monitored table changes.
- All tables updated – triggers only after all selected tables update .

Step-4

Advanced functionality(Optional)

- Min. time between triggers - Prevents over-triggering by enforcing a mandatory gap between consecutive job runs
- Wait after last change - Delays job execution until all data updates have landed in source table

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Creating and Scheduling Jobs

DEMONSTRATION

Automating Workloads with Scheduling and Triggers

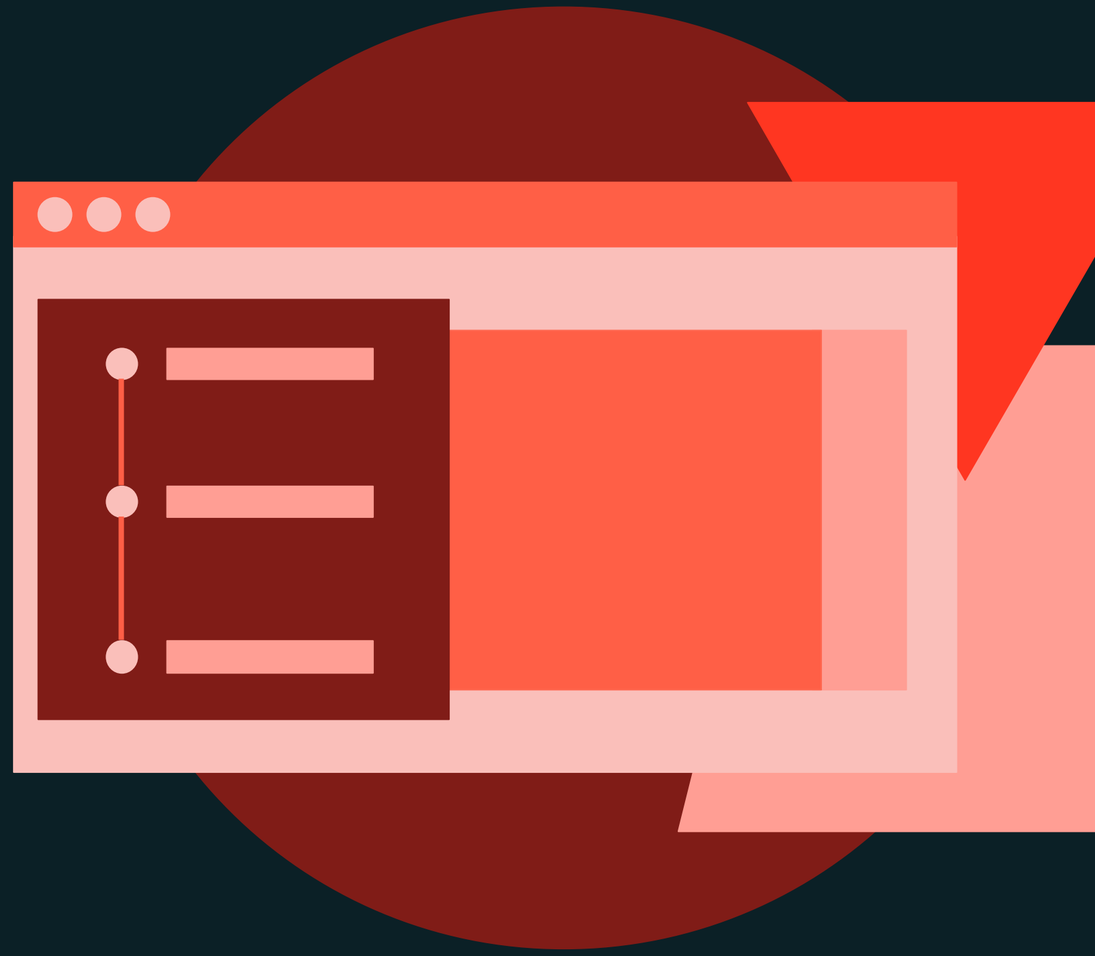

3 Demo - Automating Workloads with Scheduling and Triggers

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Advanced Lakeflow Jobs Features

Deploy Workloads with Lakeflow Jobs

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Section Learning Objectives

- Implement conditional task execution using run-if dependencies and If/Else logic.
- Utilize the repair run feature to efficiently recover from failed job runs.
- Design production-ready architectures with modular orchestration patterns, Git integration, and appropriate compute selection (serverless vs interactive clusters) for scalable workload management.
- Optimize performance monitoring and cost management using system tables (system.lakeflow) and Spark UI insights to identify bottlenecks, track SLAs, and implement resource efficiency measures.
- Apply enterprise best practices including service principal usage, parameterized tasks, alert configuration, and maintainable job design patterns to ensure reliable production deployments.

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Agenda

Advanced Lakeflow Jobs Features
• Conditional and Iterative Tasks
• Handling Task Failures
• Lakeflow Jobs in Production
• Best Practices

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Advanced Lakeflow Jobs Features

LECTURE

Conditional and Iterative Tasks

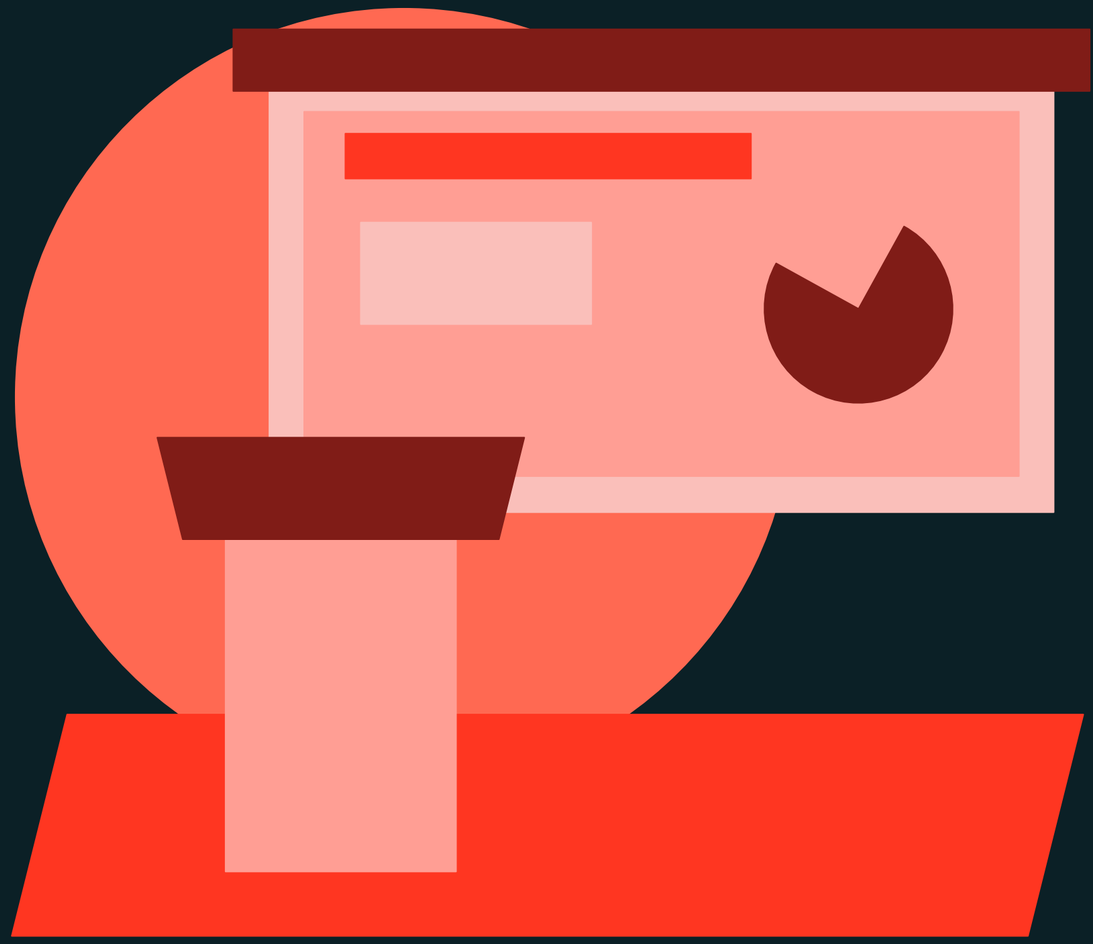

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Conditional and Iterative Tasks

Overview

Run if Conditional Task Dependencies

If Else Task

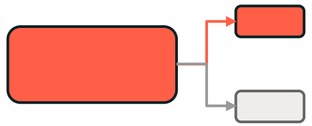

If/else

For Each Task

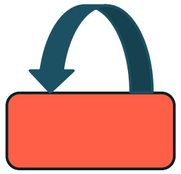

For each

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Conditional and Iterative Tasks

Run if Conditional Task Dependencies

Run if Conditional Task Dependencies

- Control task execution based on outcomes of **upstream tasks**
- Supports a variety of **dependency conditions** like:
  - All succeeded
  - At least one succeeded
  - None failed
  - And more

Example - At Least One Succeeded

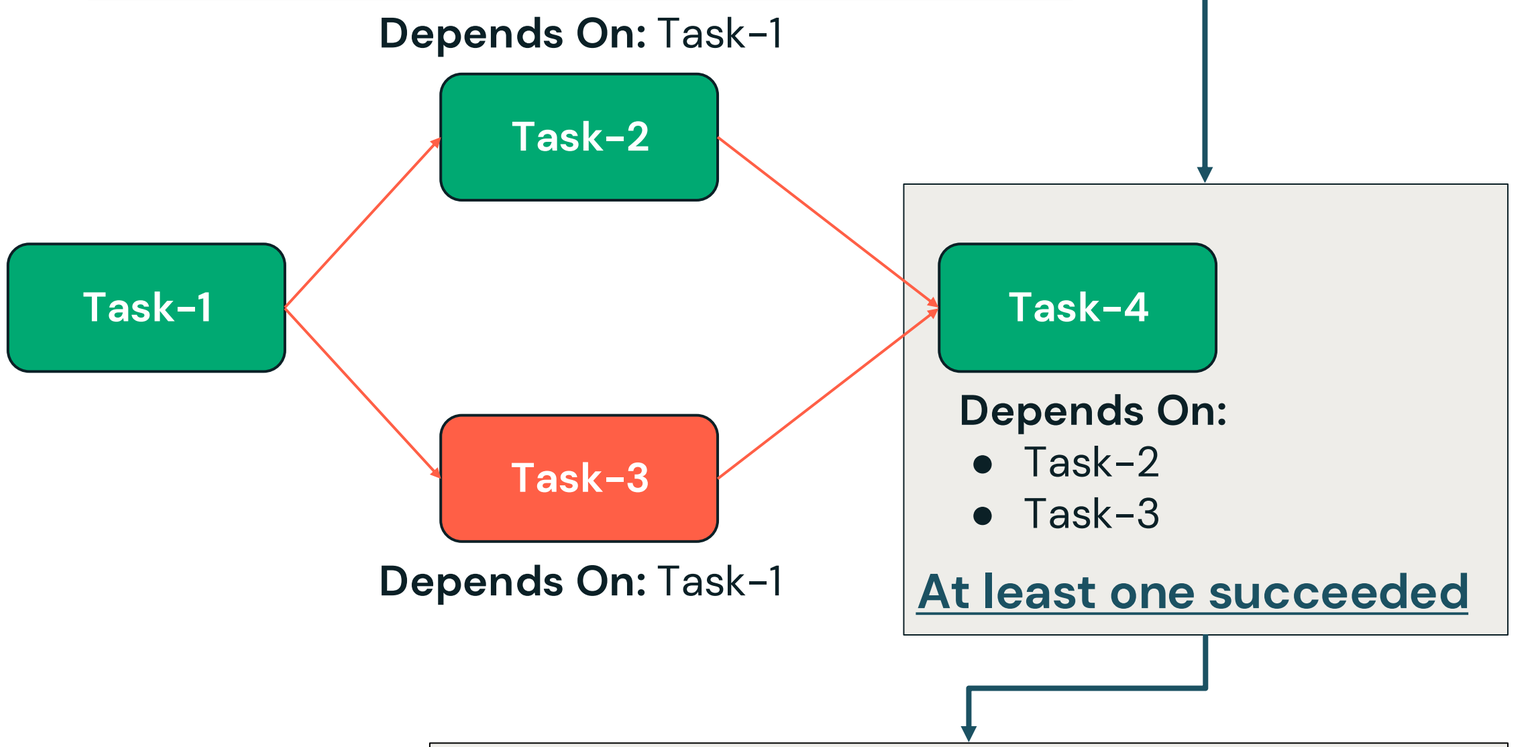

In this scenario, even if Task-3 Fails, Task-4 still runs

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Conditional and Iterative Tasks

Run if Conditional Task Dependencies

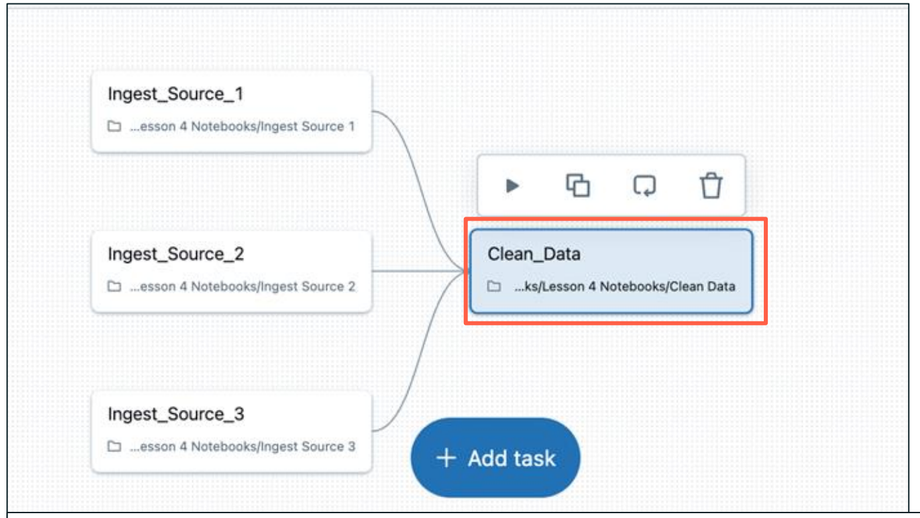

- Task dependencies provide **control flow logic based** on upstream task status
- Visually represented in the job DAG

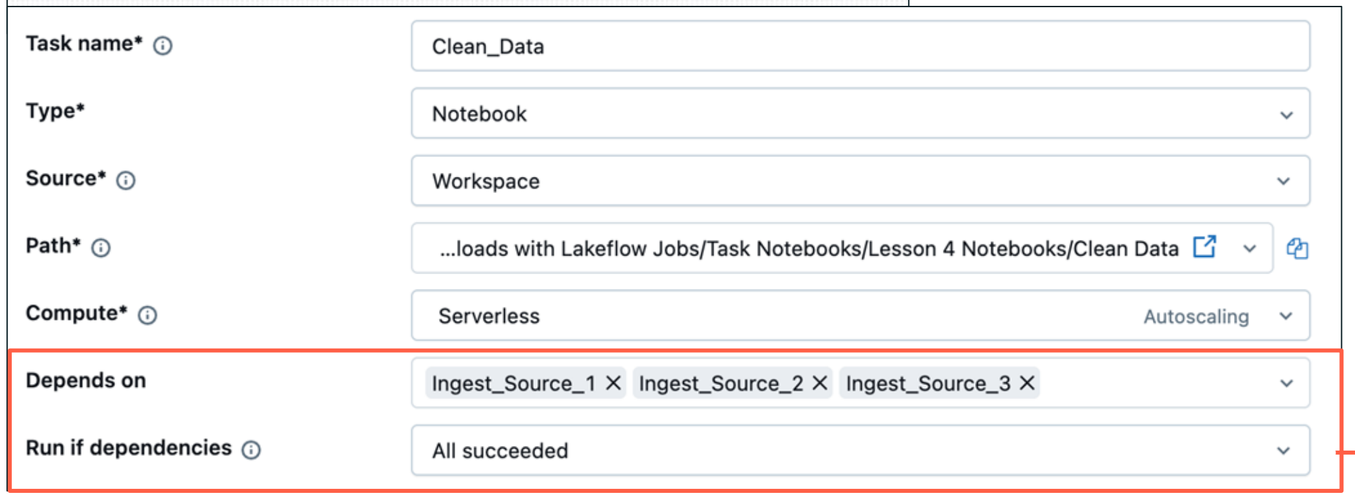

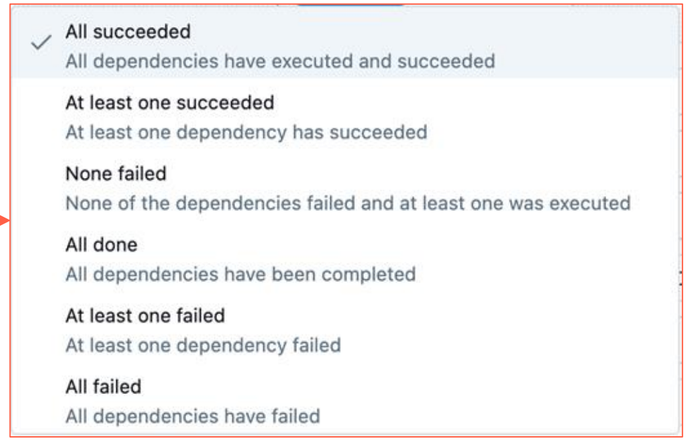

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Conditional and Iterative Tasks

Run if Conditional Task Dependencies

1. **Job Initialization**

The first three tasks are executed

2. Waiting for all task to complete (*All succeeded*)

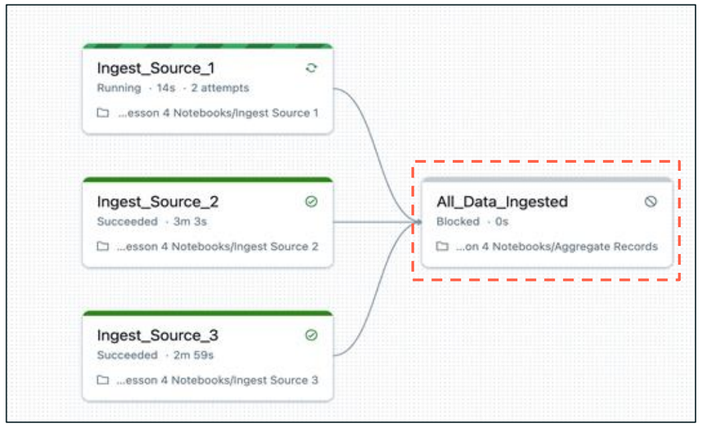

3. Dependent Task Started

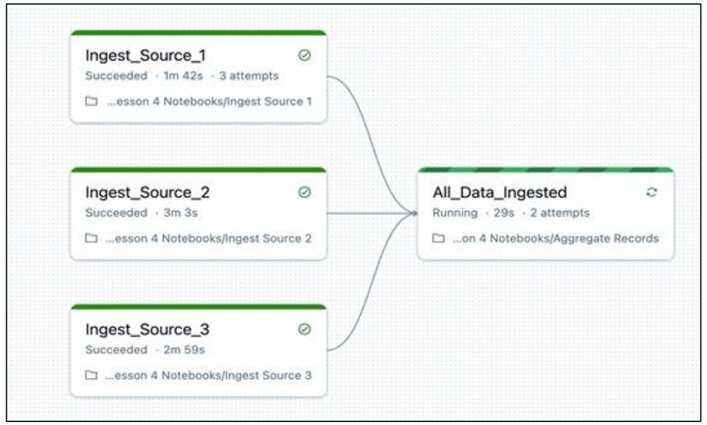

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Conditional and Iterative Tasks

If Else Task

Run if Conditional Task Dependencies

If Else Task

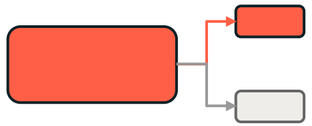

If/else

For Each Task

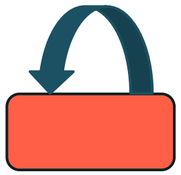

For each

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Conditional and Iterative Tasks

If/else Conditional Tasks

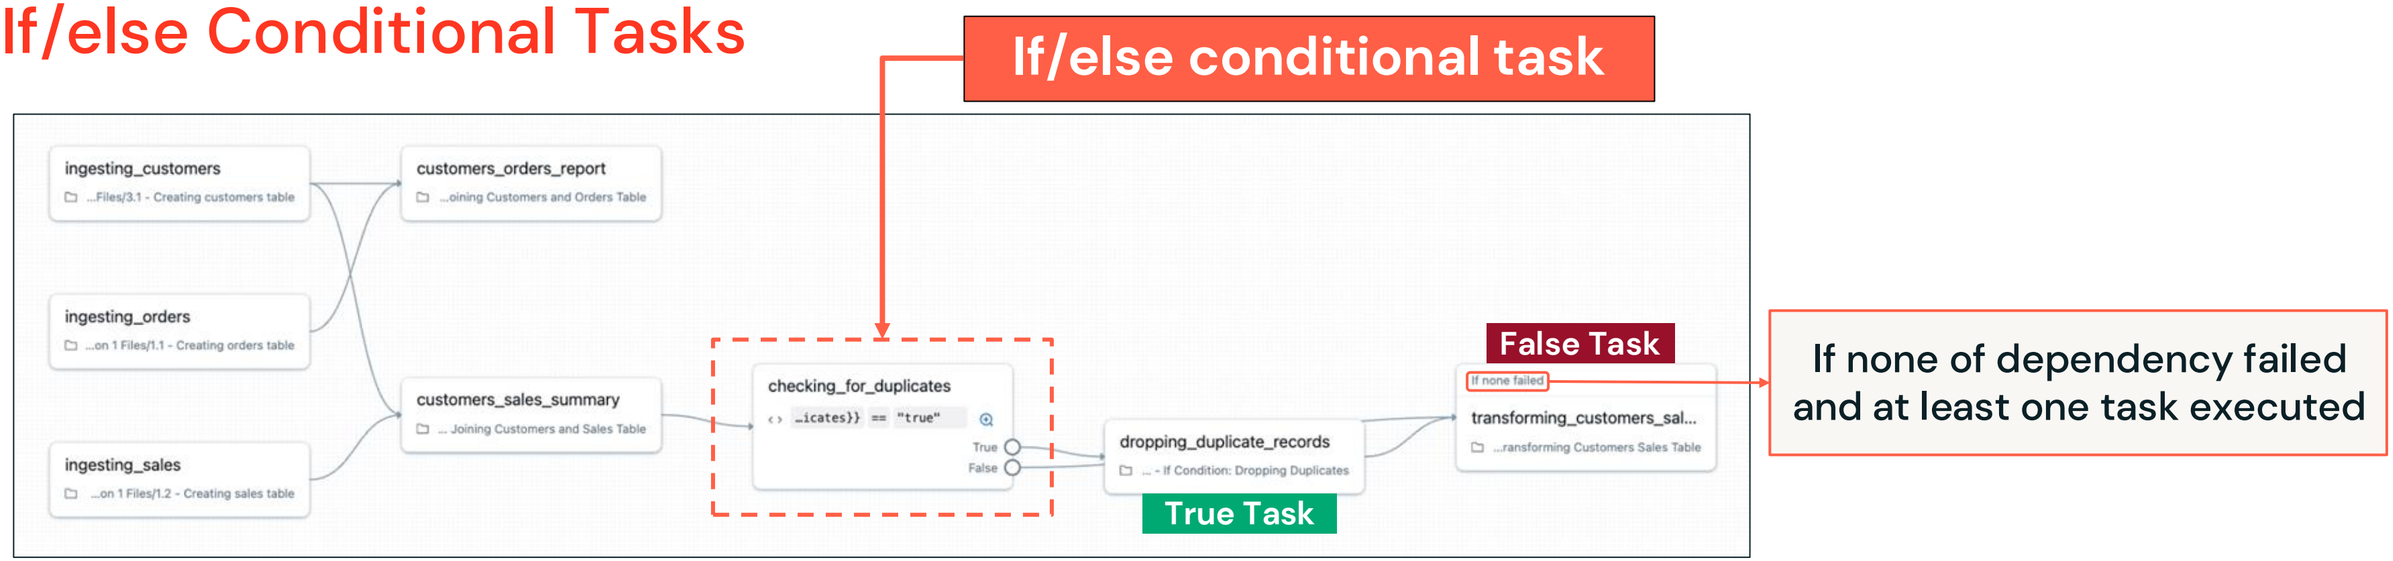

If/else Conditional Tasks

- Adds boolean conditional logic to your workflow based on task results
- Allows branching based on specific conditions (e.g., data quality checks, record counts)
- Uses boolean operators (==, !=, >, >=, <, <=) to evaluate expressions

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Conditional and Iterative Tasks

For Each Task

Run if Conditional Task Dependencies

If Else Task

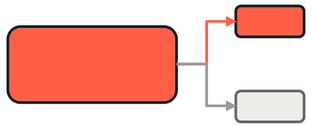

If/else

For Each Task

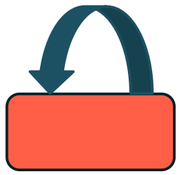

For each

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Conditional and Iterative Tasks

For Each Task

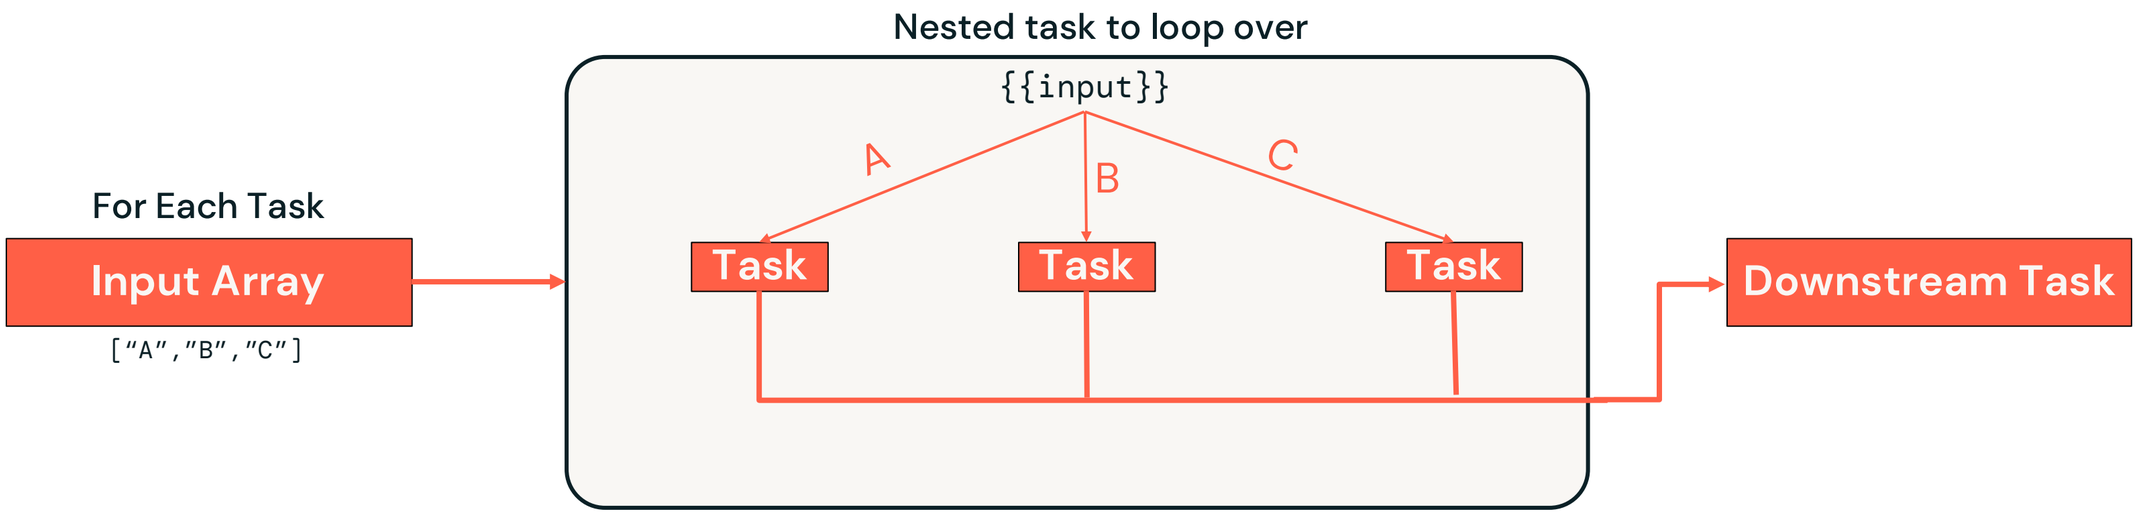

For Each Task

- For Each loops over an input array and runs the same nested task per item, passing it as `{{input}}`
- Iterations can run in parallel with configurable concurrency to speed up execution.
- Downstream dependencies attach to the For Each container, not to the nested task.

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Conditional and Iterative Tasks

For Each Task

1. For Each Task

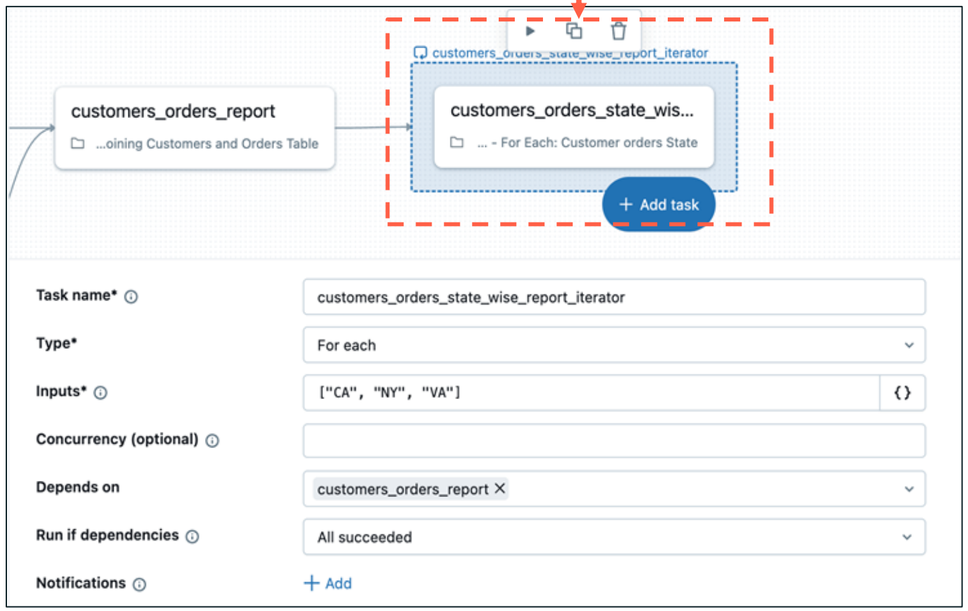

2. Nested Task

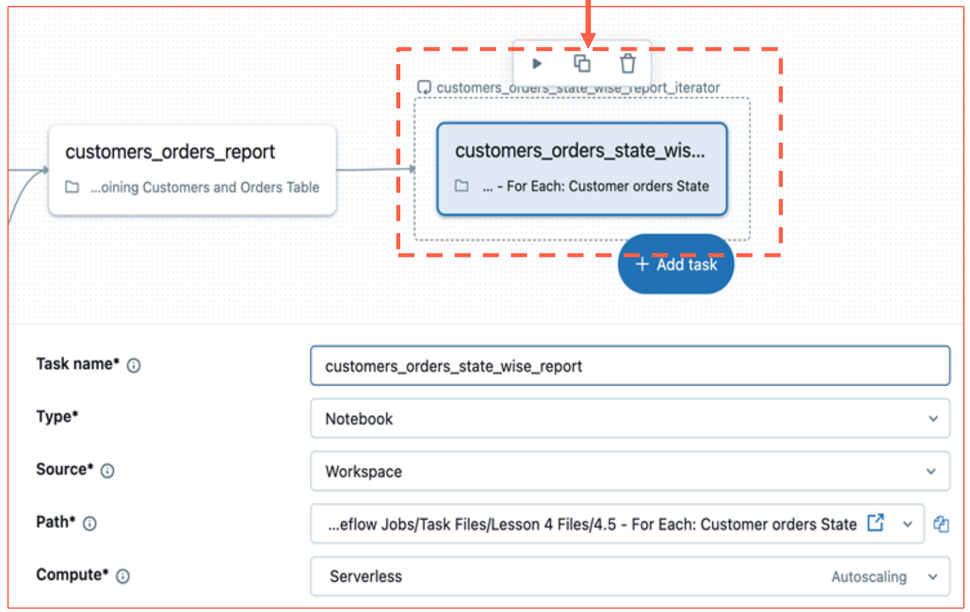

For Each task requires defining two components:

1. **The For Each task –** The top-level container task that manages the loop
2. **A nested task –** The actual task that executes for each iteration

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Advanced Lakeflow Jobs Features

DEMONSTRATION

Building Dynamic Workloads with Advanced Tasks

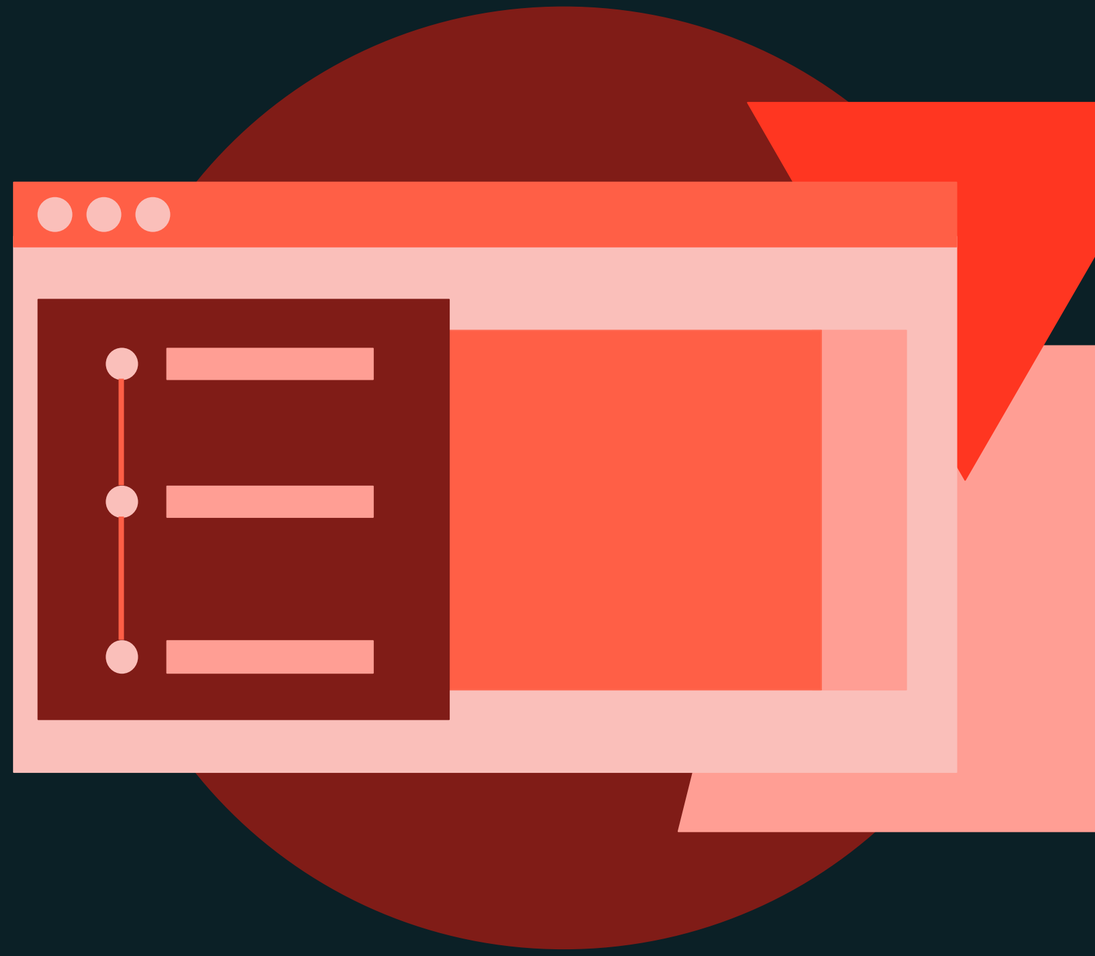

4 Demo - Building Dynamic Workloads with Advanced Tasks

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Advanced Lakeflow Jobs Features

LAB EXERCISE

Adding If-Else Task and Automating your Job

5 Lab - Adding If-Else Task and Automating your Job

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Advanced Lakeflow Jobs Features

LECTURE

Handling Task Failures

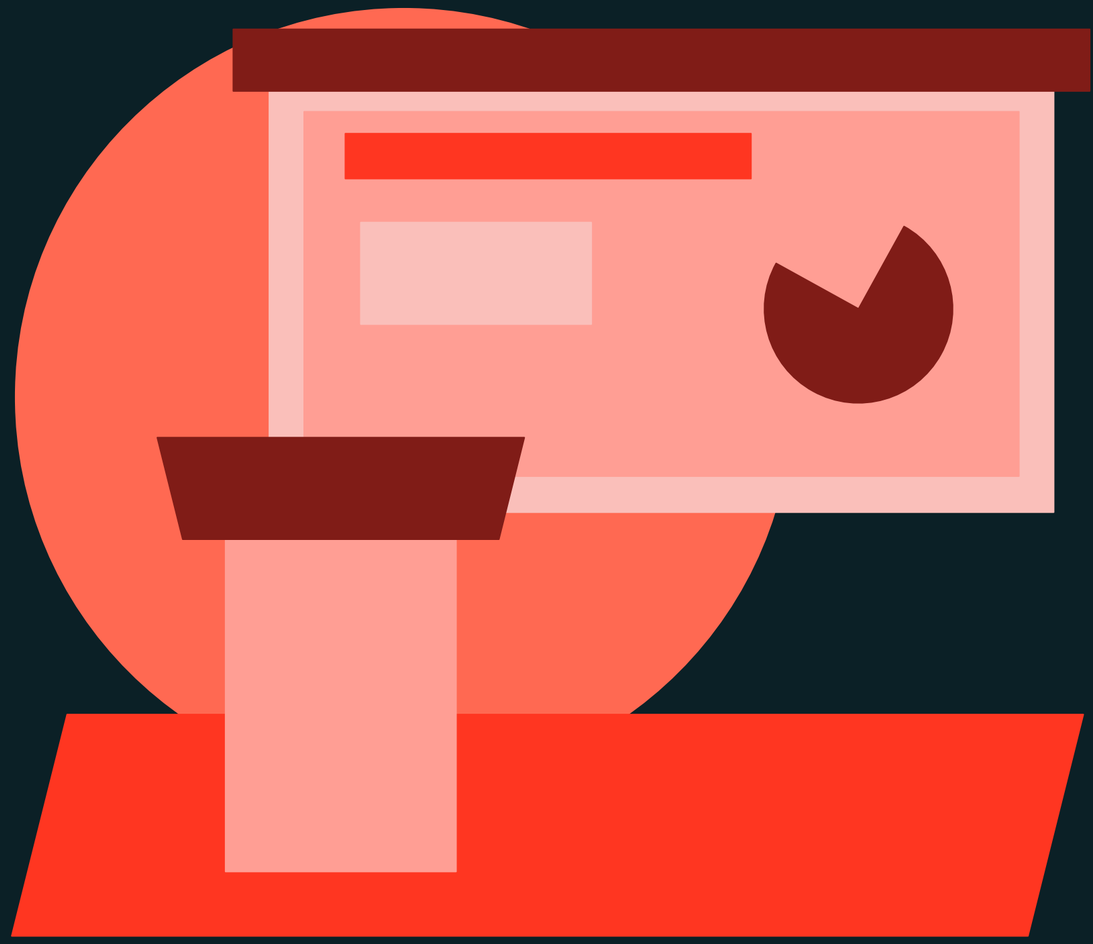

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Handling Task Failures

Repair and Rerun

Re-run the **failed** task

- **Repair** feature allows to **re-run** and override task parameters
- **Reduces the time** and resources required to recover from unsuccessful job runs
- In case of **Task Failure,** You can:
  - Modify the **task** and run again
  - Modify the **parameters** and run again

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Handling Task Failures

Repair and Rerun

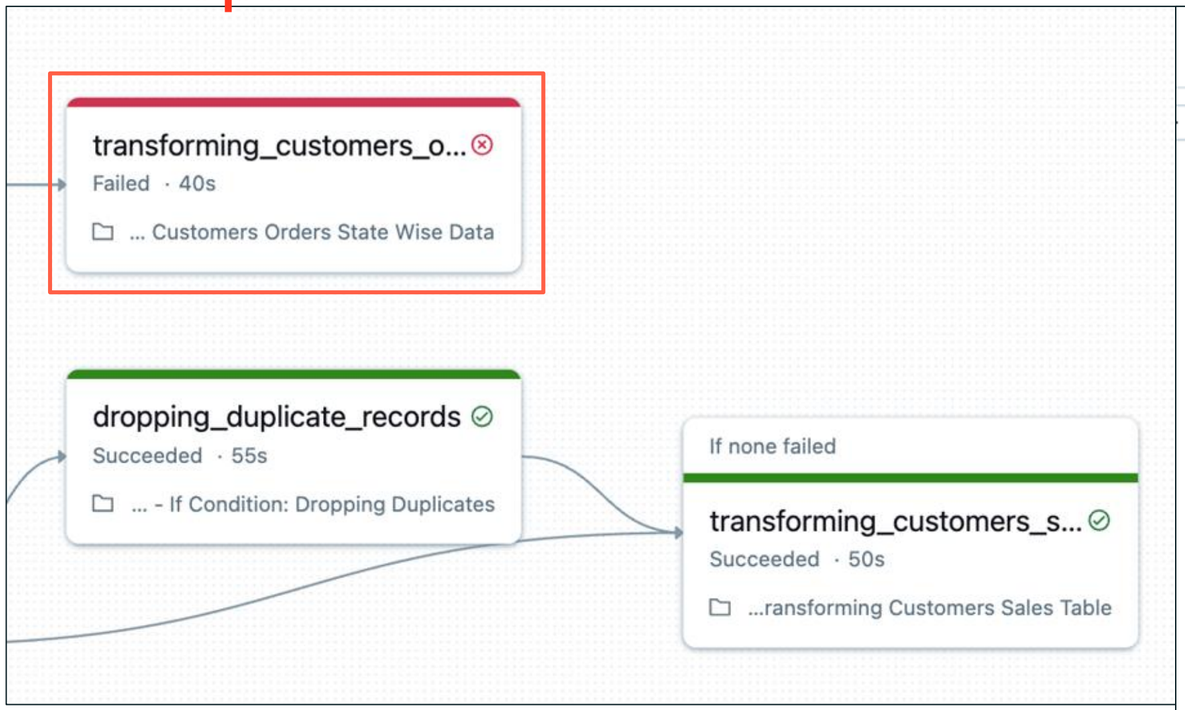

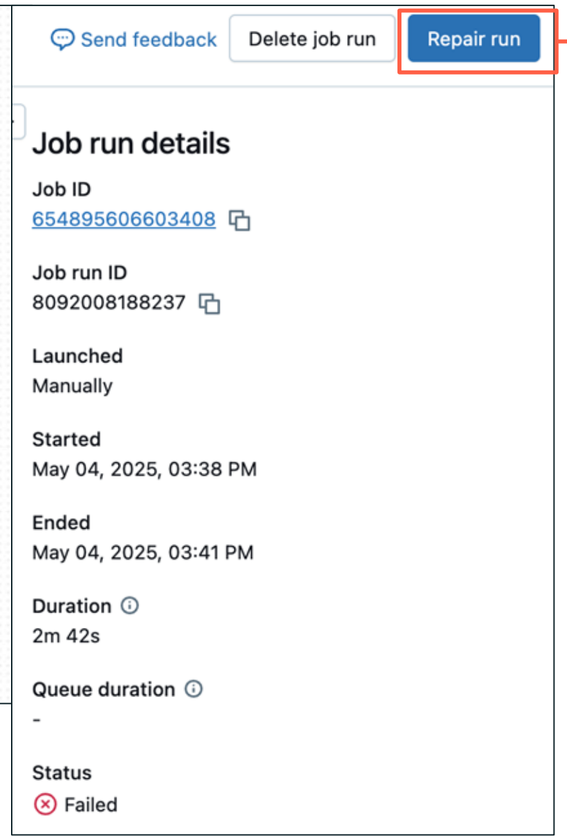

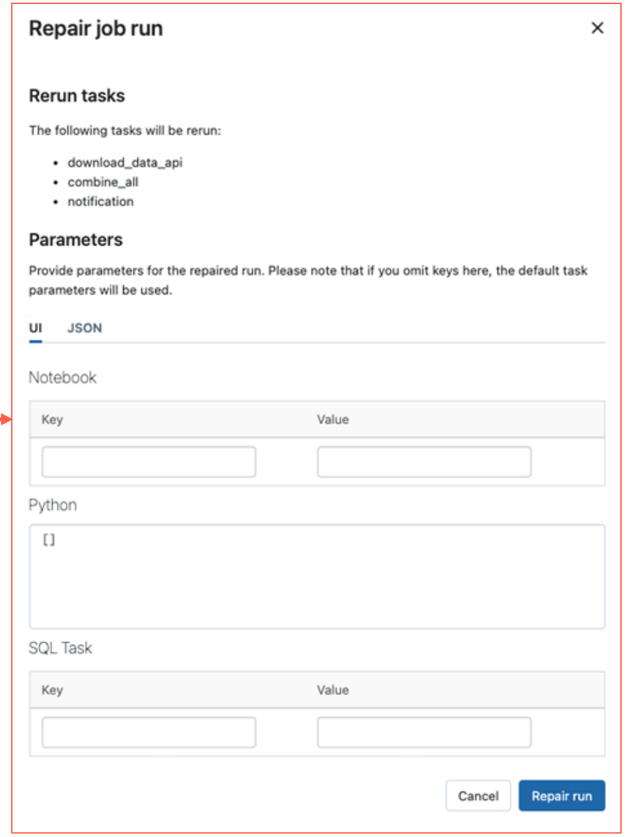

Allowing you to **run only failed tasks**, saving you time and money by only re-running the necessary tasks instead of the entire job

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Handling Task Failures

Repair and Rerun

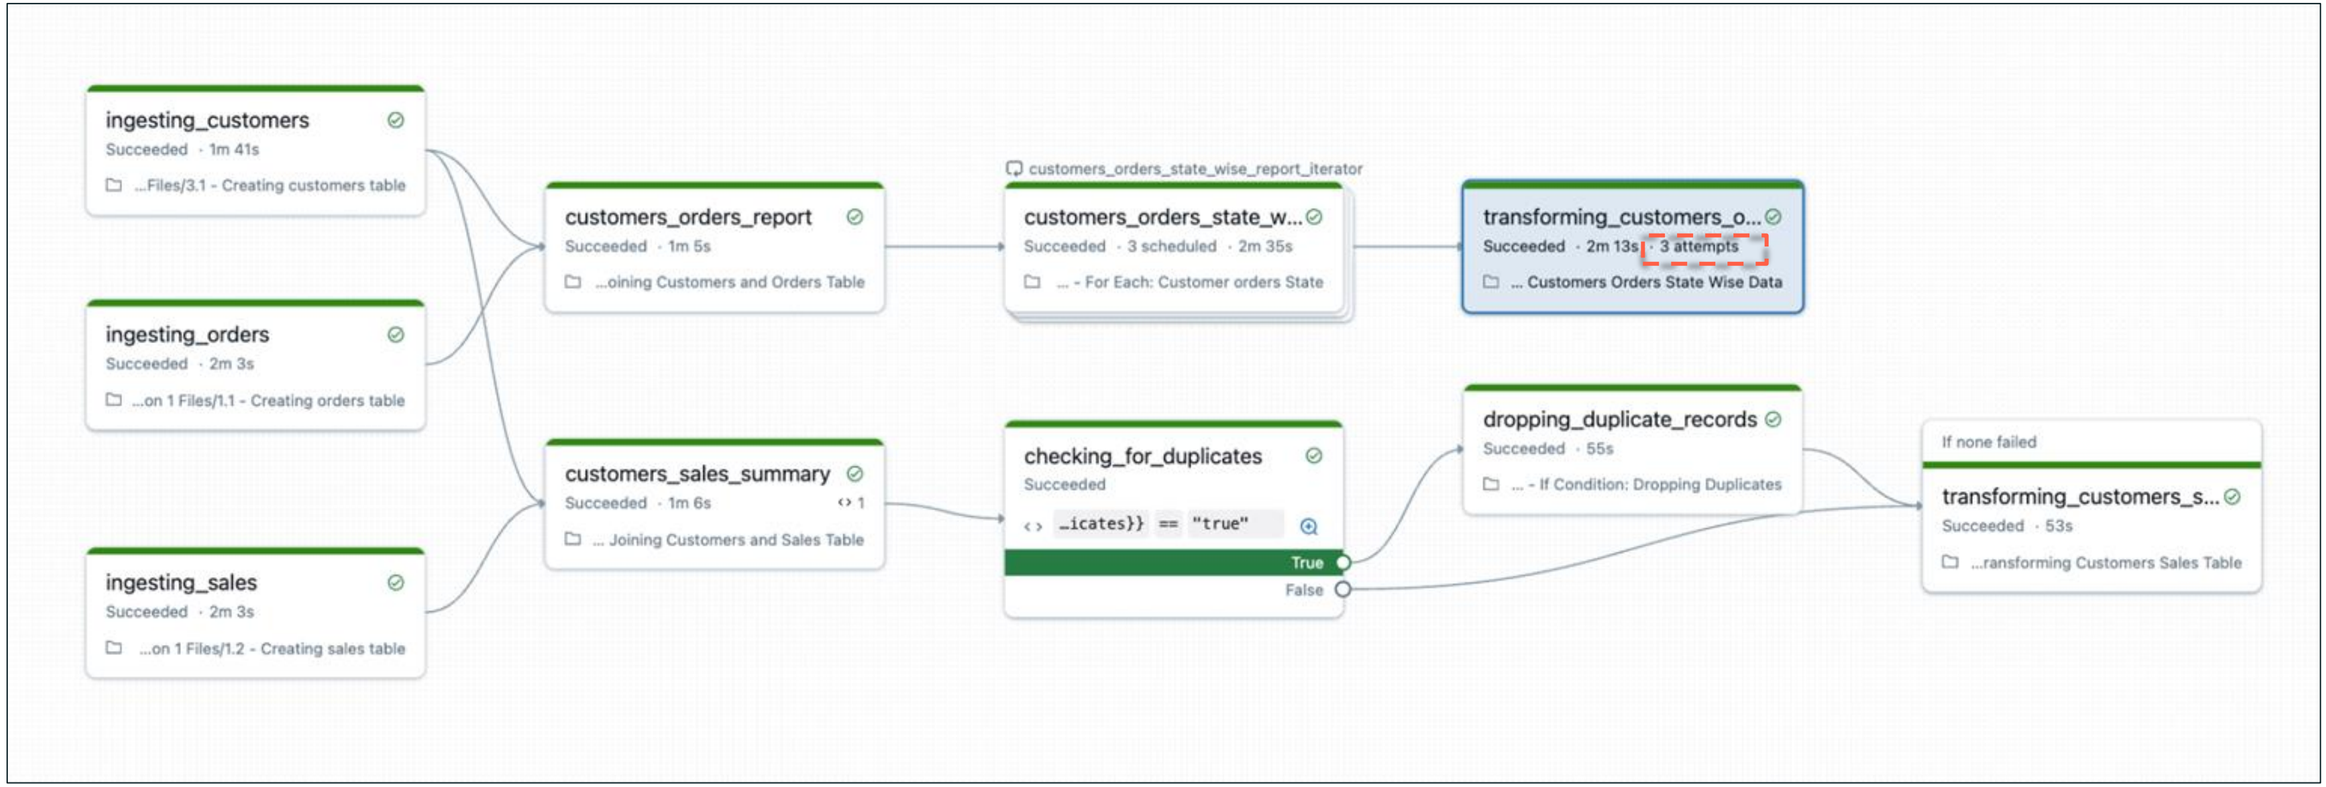

After re-running task, your final job will look like this.

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Advanced Lakeflow Jobs Features

LECTURE

Monitoring Jobs Performance

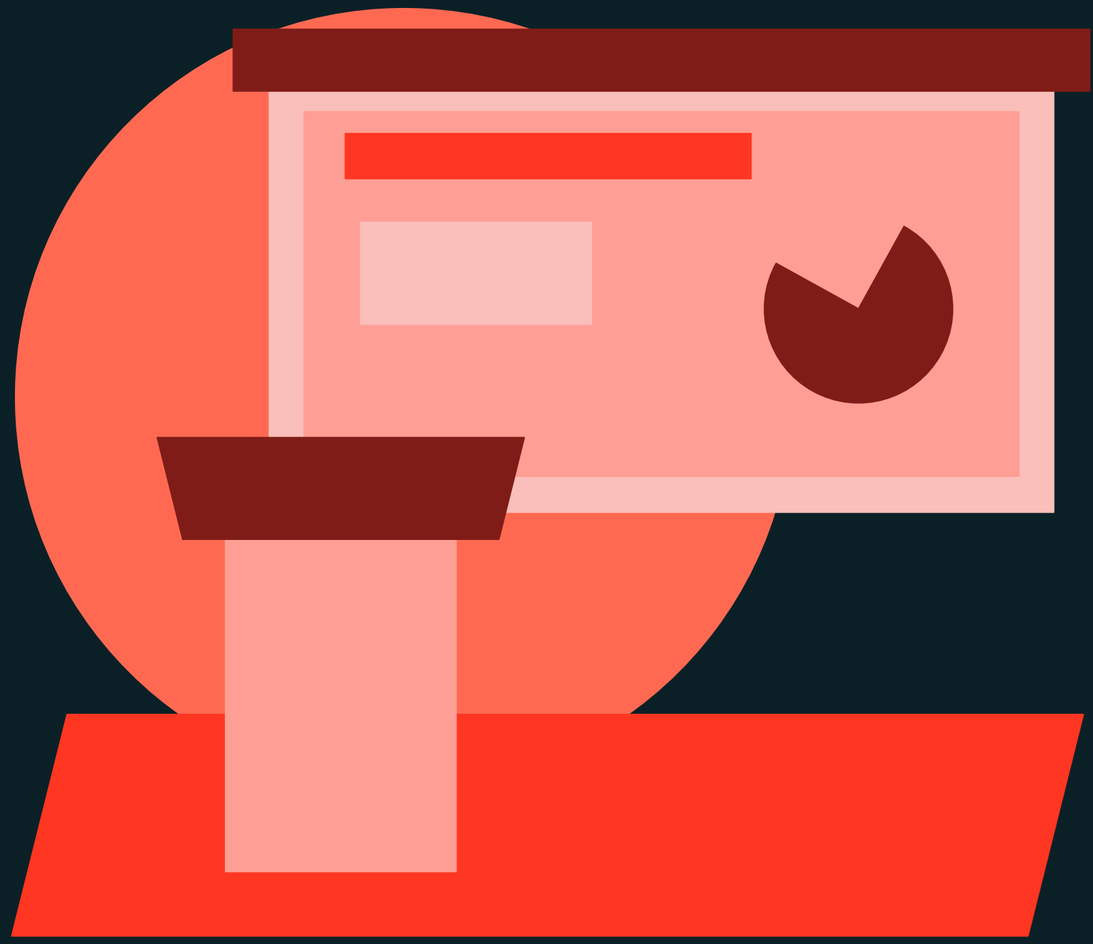

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Monitoring Jobs Performance

System Tables

- **system.lakeflow** is a built-in, read-only catalog that logs all job activity across workspaces in the region
- Timeline tables slice long runs hourly using **period_start_time** and **period_end_time**, enabling reliable duration, concurrency, and SLA analytics.
- Key Tables:
  - **jobs** (Job basic info)
  - **job_tasks** (Task basic definitions)
  - **job_run_timeline** (Each job run over time)
  - **job_task_run_timeline** (Each task run over time)
  - **Pipelines** (Pipelines basic info)

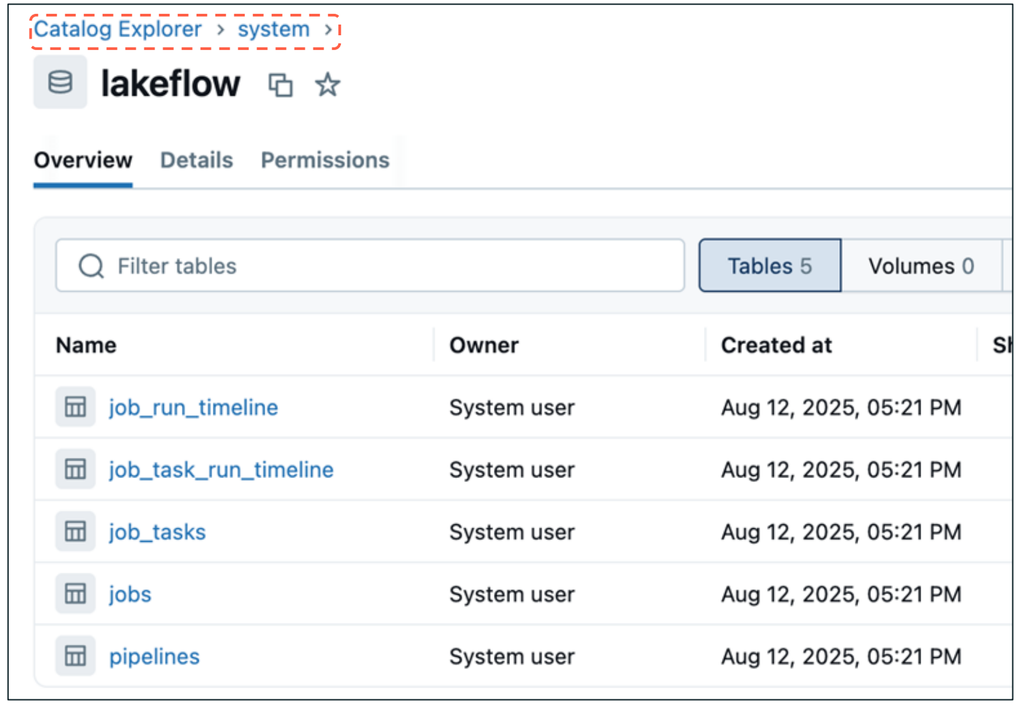

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Monitoring Jobs Performance

Spark UI

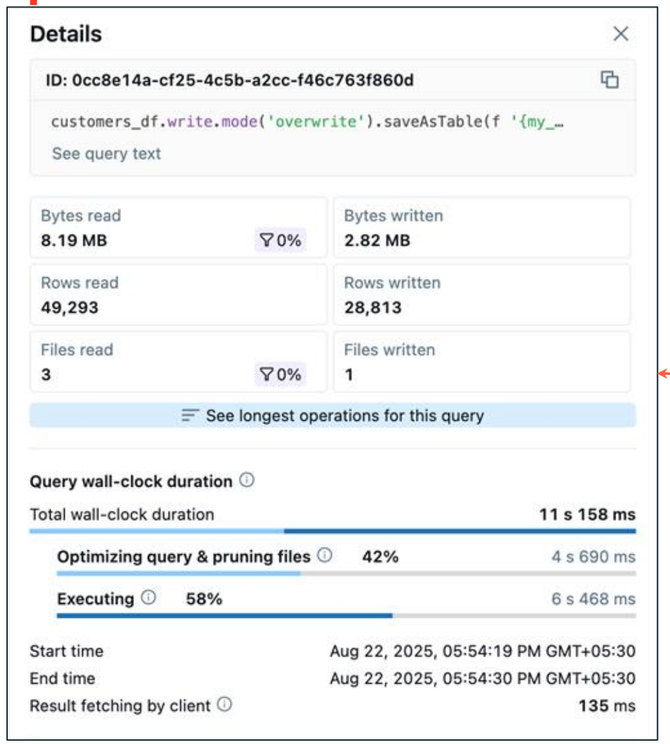

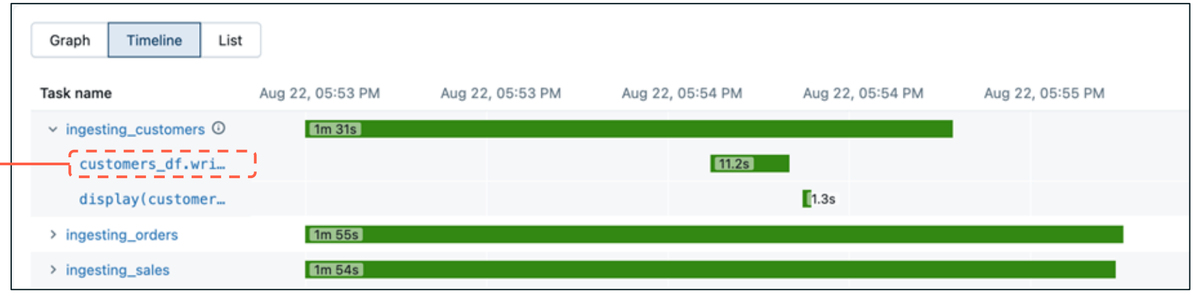

- **Timeline** in job run highlights task start/end, duration, and overlap to spot bottlenecks fast.
- Click a task to see status, timestamps, duration, cluster/runtime, logs, and quick I/O.
- Click on **query/code** details for full text, run ID, wall-clock split (optimizing/pruning vs executing), files read/written details, files & partitions and spill details

Act on insights:
High planning time → Improve pruning/partitioning
High execution time → Optimize joins/aggregations (broadcast, skew fixes)

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Advanced Lakeflow Jobs Features

DEMONSTRATION

Monitoring and Repairing Task

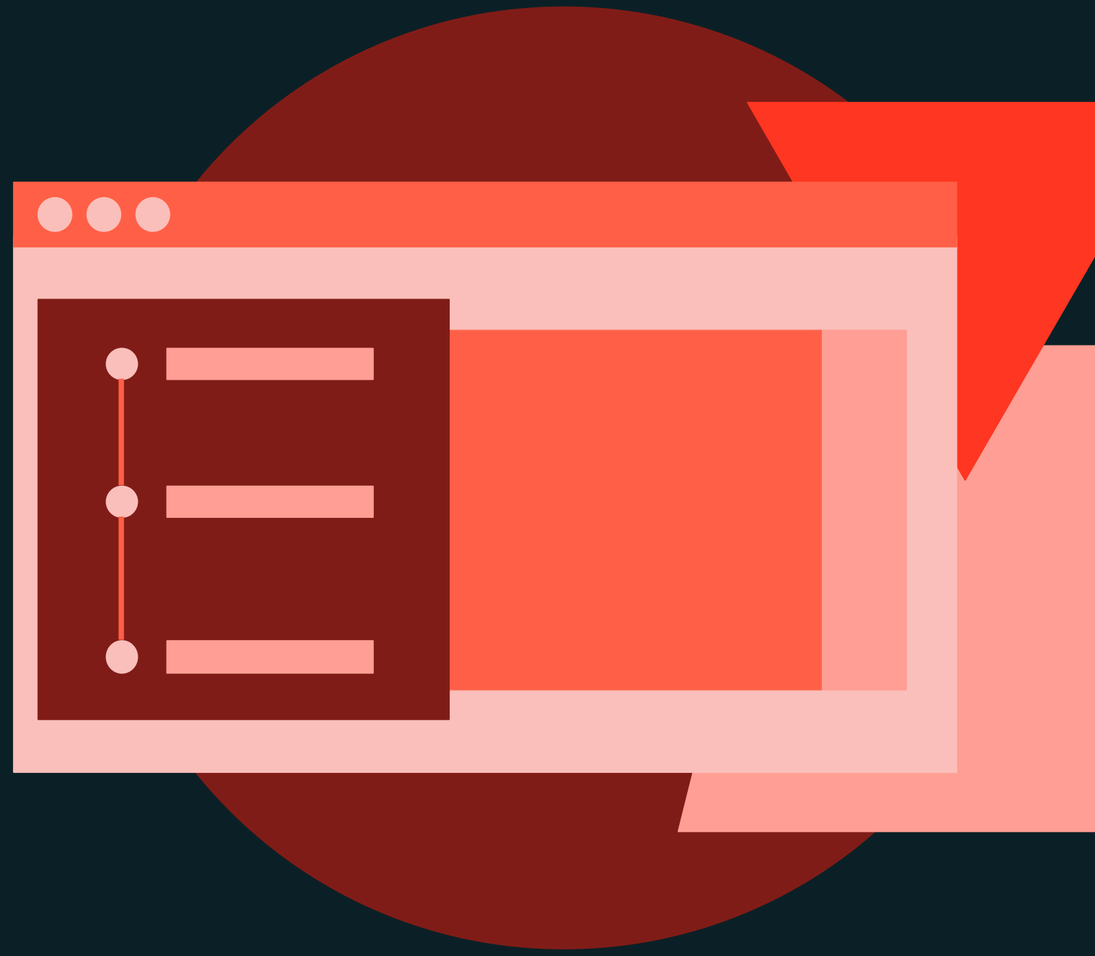

Notebook: 6 – Monitoring and Repairing Task

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Advanced Lakeflow Jobs Features

LECTURE

Lakeflow Jobs in Production

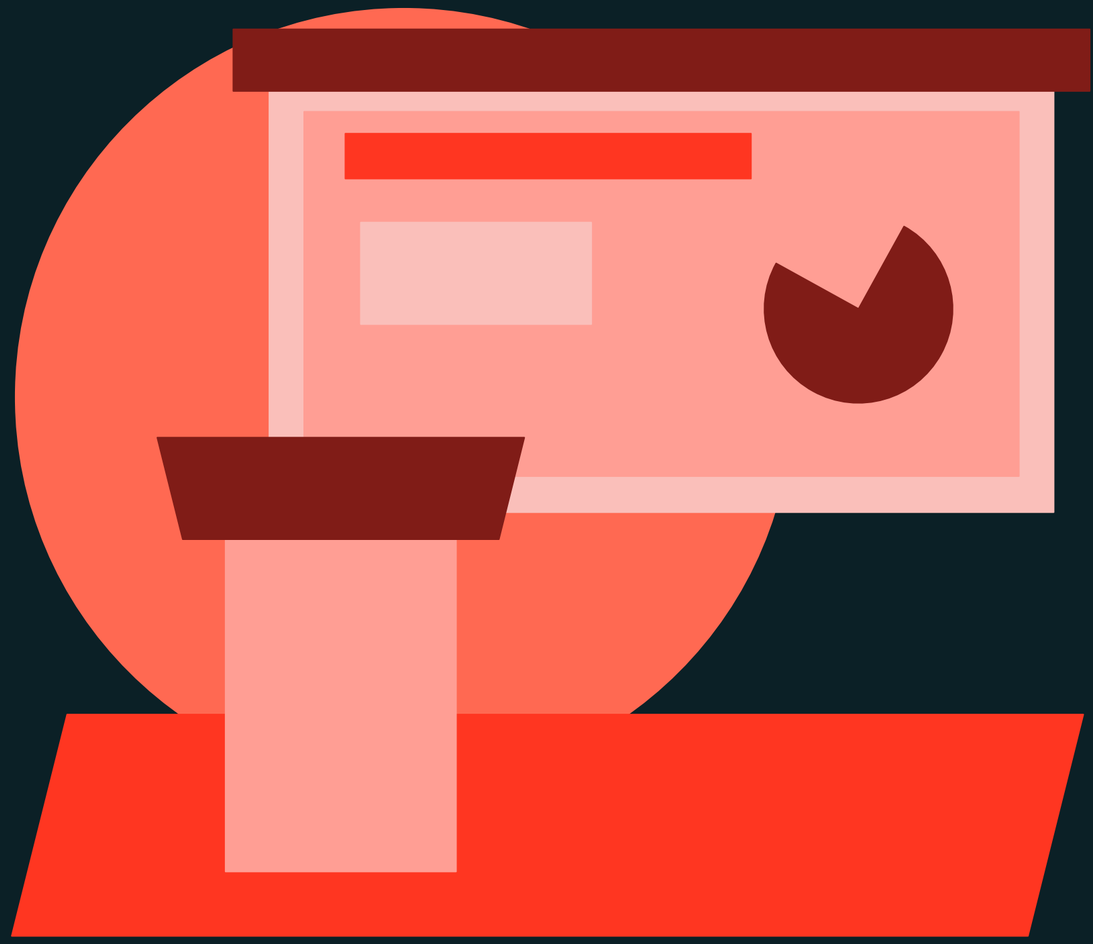

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lakeflow Jobs in Production

Common Best Practices

Compute

Modular Design

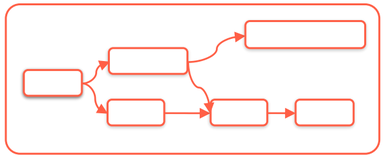

Git

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lakeflow Jobs in Production

Selecting Compute

Interactive Clusters

It is best for performing ad-hoc analysis, data exploration, or development, **but not production**

- Costly for Job runs
- Limited Scalability
- Availability could be an issue because of parallel usage

Job Clusters

It is **cheaper** as they terminate when the job ends, reducing resource usage and costs. But,

- Start-up Latency
- Subject to cloud provider start-up time
- Maintenance Burden

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lakeflow Jobs in Production

Selecting Compute

Serverless Compute

Simplicity: You no longer need to choose the VM type. Customers do not need dedicated skilled DevOps folks

High efficiency: Photon turned on default

Faster startup: VMs run in Databricks account. We ensure enough nodes are available based on historical data and ML

Reliability: Shielded from cloud disruptions

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lakeflow Jobs in Production

Pricing Structure

Classic

DBUs
(paid to Databricks)

Infrastructure cost
(paid to cloud provider)

- Cost of VMs for clusters
- Cost of network (FW / NAT)
- Cost of security, utilization monitoring,

Operational cost
(incurred at organization level)

- Time spent on deploying, automating, maintaining infrastructure
- Time spent on managing cost, efficiency, utilization

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lakeflow Jobs in Production

Pricing Structure

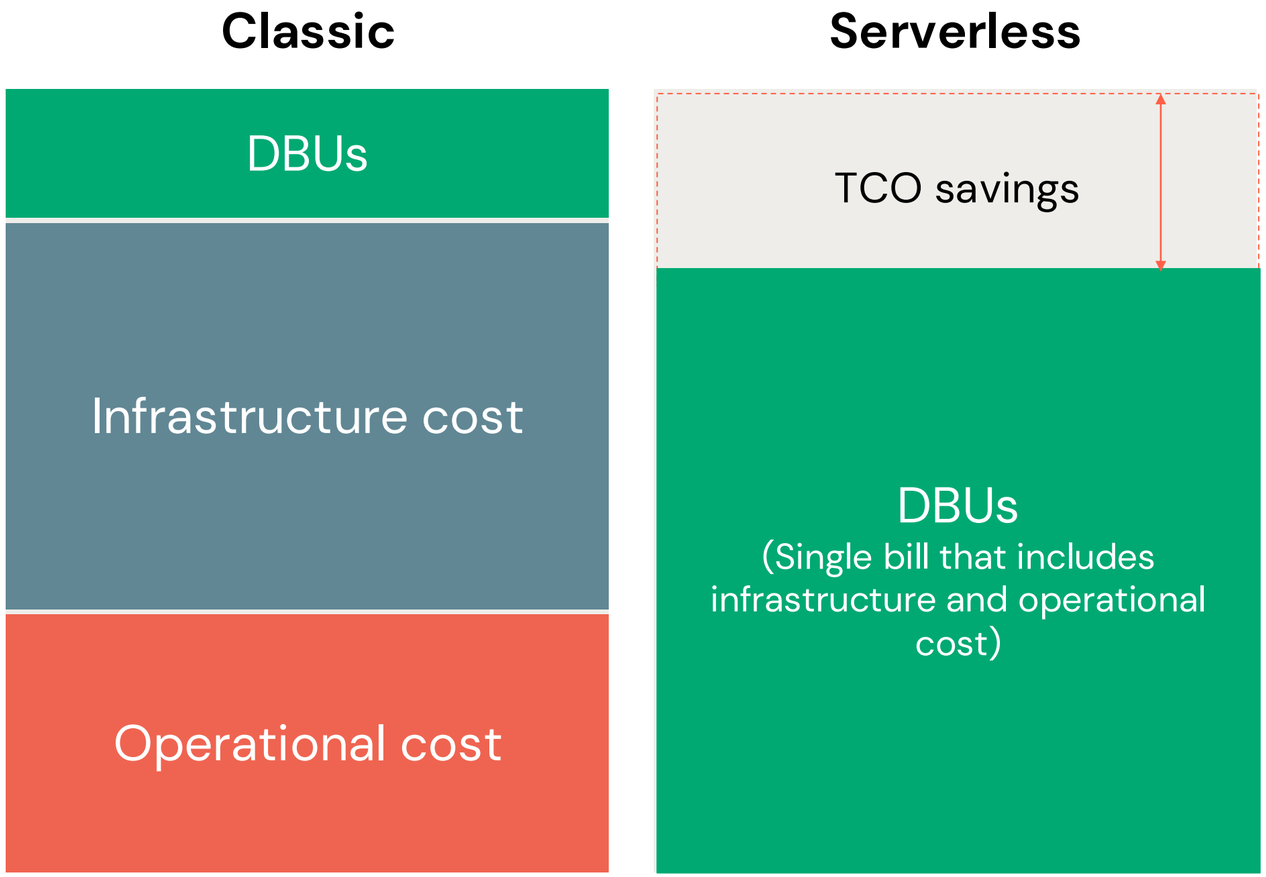

Value

Fully managed service - operationally simpler, more reliable

Fast clusters, auto-scaling - better user experience, lower cost

Out of the box performance and optimizations - lower overall TCO

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lakeflow Jobs in Production

Common Best Practices

Compute

Modular Design

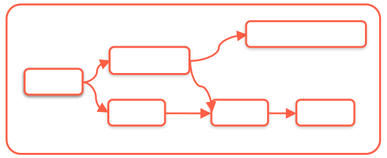

Git

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lakeflow Jobs in Production

Modular Design in Databricks LakeFlow Jobs

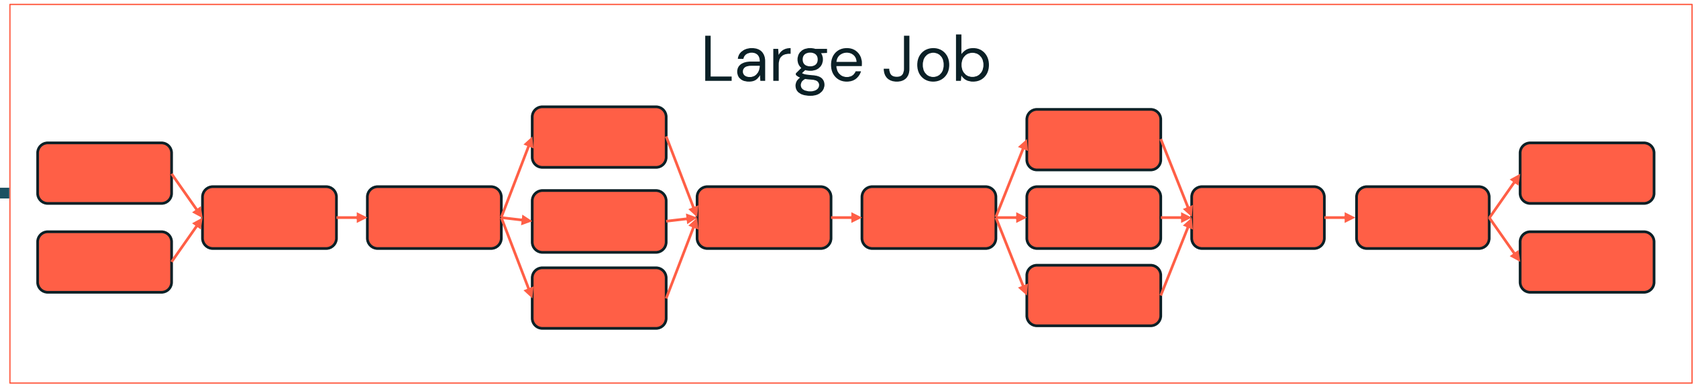

Modular orchestration in jobs enables you to break down complex DAGs into smaller, manageable components

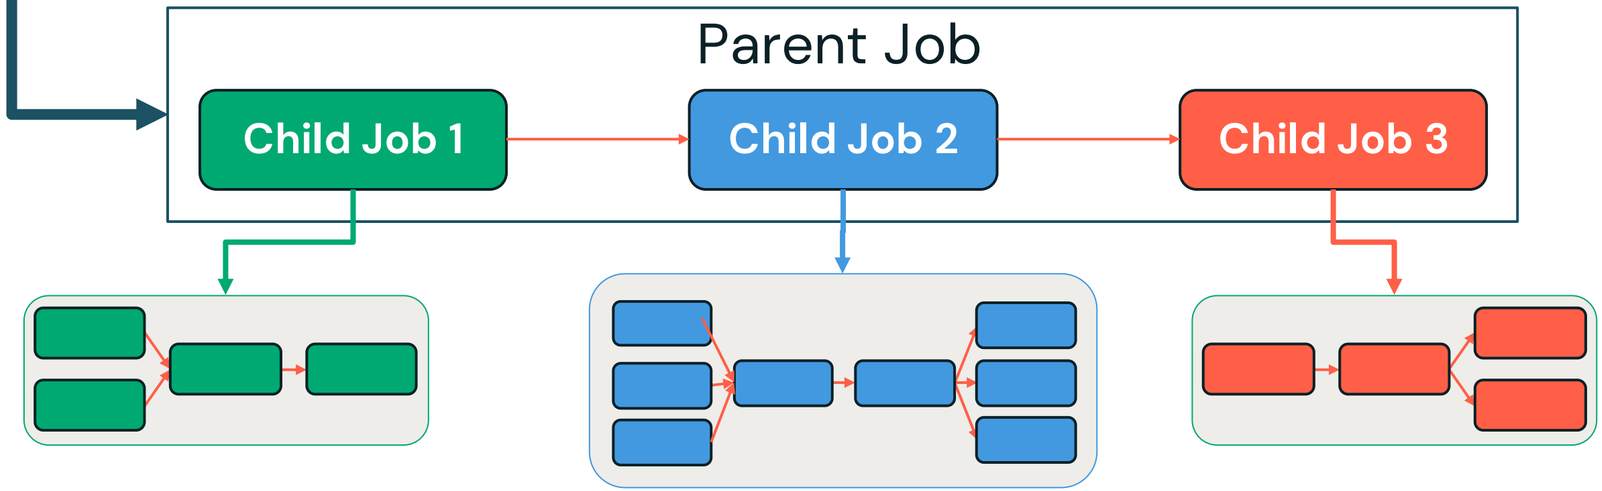

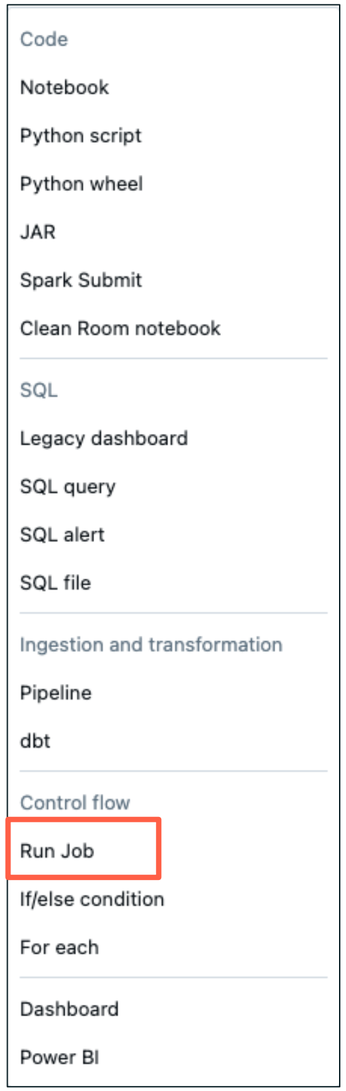

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lakeflow Jobs in Production

Common Best Practices

Compute

Modular Design

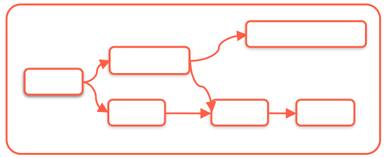

Git

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lakeflow Jobs in Production

Jobs and Git

Jobs supports execution of notebooks from Git

- This **prevents unintentional changes** to your **production** job, for example, when another user makes local edits in a 'Prod' repo or switches branches
- **Simplifying the job definition** process by having a single source of truth
- Makes **deployment** easy for notebooks through CI/CD
- Users can **connect** with **various Git platforms**, as Databricks supports GitHub, GitLab, AWS CodeCommit, and other Git providers.

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Lakeflow Jobs in Production

Jobs and Git Configuration Steps

Create task with Git provider as Source

1

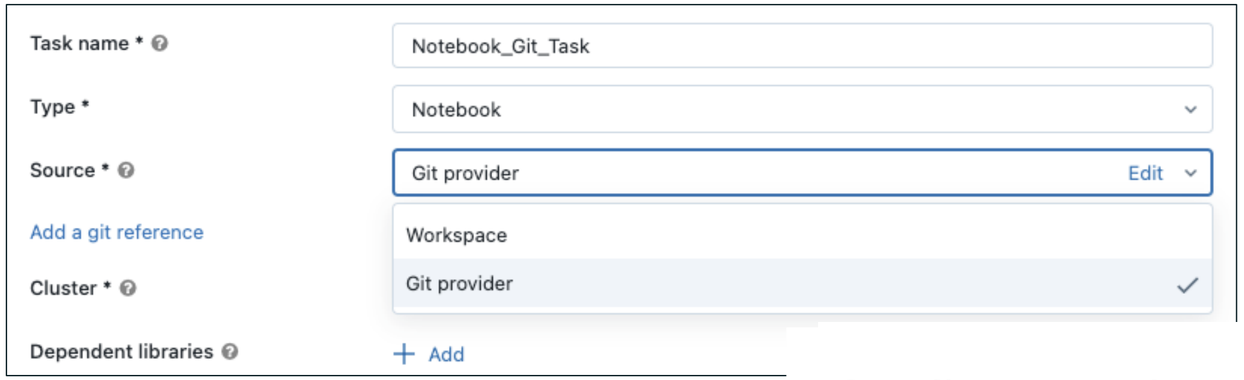

Configure path to main notebook under repository root

2

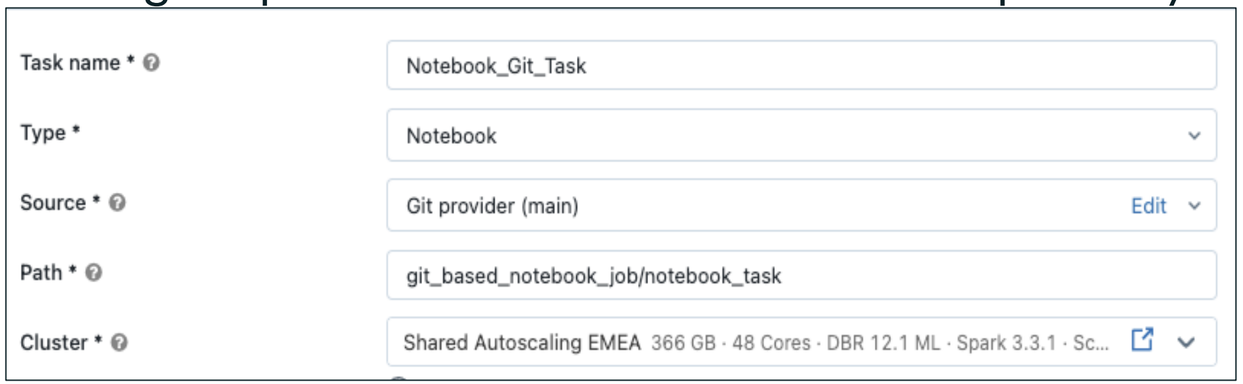

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Advanced Lakeflow Jobs Features

LECTURE

Best Practices

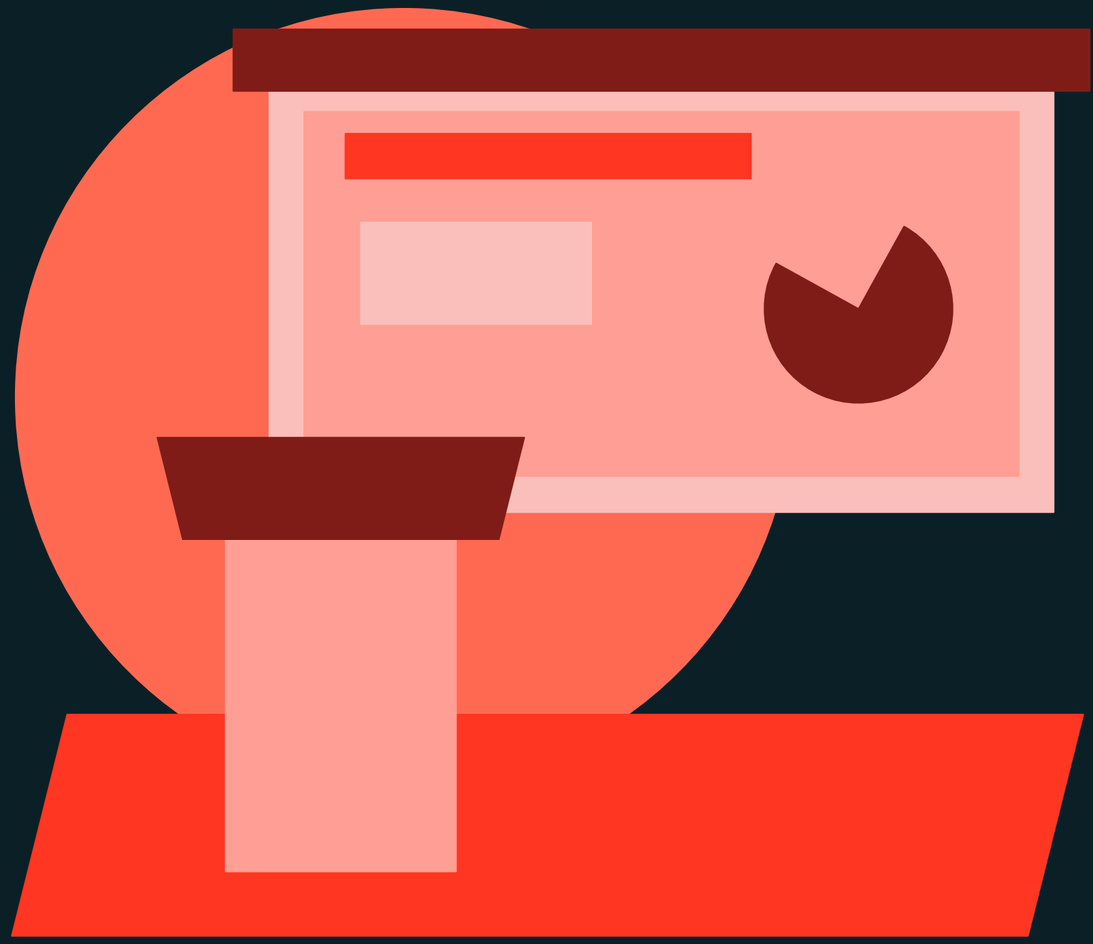

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Best Practices

Compute & Cost Optimization

- Use job or **serverless clusters** in production
- Avoid interactive clusters for production workloads
- **Enable Photon** for faster and cheaper execution
- **Reuse clusters** to reduce startup time and cost

Orchestration & Modularity

- **Break complex pipelines** into modular jobs/run jobs
- Use **multi-task jobs** for parallel and scalable execution
- Apply conditional logic using Run If, If/Else, For Each task
- Limit jobs to **manageable task counts** for maintainability (Limited to 1000 task)

Monitoring & Governance

- **Use service principals** for job ownership and data access
- Configure alerts for **failures, delays, and completions**
- Leverage **Repair & Run** for reducing cost on re-processing
- Parameterize tasks for **reusability and flexibility**

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Summary and Next Steps

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Sp
Iceberg logo are trademarks of the Apache Software Foundation.

Course Learning Objectives Recap

- Understand the role of Lakeflow Jobs within the Databricks ecosystem as a unified orchestration platform for data, analytics, and AI workloads.
- Design and implement data workloads using Directed Acyclic Graphs (DAGs), demonstrating the relationship between jobs, tasks, and dependencies.
- Configure various job scheduling options including manual, scheduled, file arrival, and continuous triggers to automate workflow execution.
- Implement advanced workflow features such as conditional task execution, run-if dependencies, and repair runs to create robust, fault-tolerant data pipelines.
- Apply best practices for production workloads including selecting appropriate compute options, implementing modular orchestration, and utilizing proper error handling techniques.

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Next Steps

Additional resources for continuing the learning journey.

Data Engineering with Databricks

- Continue your learning through self-paced or instructor-led offerings
- The courses offer hands-on instruction in:
  - Databricks Data Science & Engineering Workspace
  - Databricks SQL
  - Databricks Repos
  - Databricks Task Orchestration
  - Unity Catalog

Data Engineer Associate Certification

- Validate your data and AI skills on Databricks by earning a Databricks credential
- Exam information and exam guide
- The exam covers:
  - Data Intelligence Platform
  - ELT With Spark SQL and Python
  - Incremental Data Processing
  - Production Pipelines
  - Data Governance

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

Advanced Lakeflow Jobs Features

LAB EXERCISE

BONUS - Modular Orchestration

This is an optional bonus lab that can be completed after class if you're interested in practicing modular design.

You will most likely not have time to complete this bonus lab during a live class.

Notebook: 7 BONUS LAB - Modular Orchestration

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.

databricks

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the Apache Software Foundation.